

Lecture 4

Instruction Level Parallelism (2)

EEC 171 Parallel Architectures

John Owens

UC Davis

Credits

- © John Owens / UC Davis 2007–17.
- Thanks to many sources for slide material: Computer Organization and Design (Patterson & Hennessy) © 2005, Computer Architecture (Hennessy & Patterson) © 2007, Inside the Machine (Jon Stokes) © 2007, © Dan Connors / University of Colorado 2007, © Wen-Mei Hwu/David Kirk, University of Illinois 2007, © David Patterson / UCB 2003–6, © John Lazzaro / UCB 2006, © Mary Jane Irwin / Penn State 2005, © John Kubiawicz / UCB 2002, © Krste Asinovic/Arvind / MIT 2002, © Morgan Kaufmann Publishers 1998.

Today's Goals

- Out-of-order execution
- An alternate approach to machine parallelism: software scheduling & VLIW

Instruction Issue and Completion Policies

- Instruction-issue—initiate execution
 - Instruction lookahead capability—fetch, decode and issue instructions beyond the current instruction
- Instruction-completion—complete execution
 - Processor lookahead capability—complete issued instructions beyond the current instruction
- Instruction-commit—write back results to the RegFile or D\$ (i.e., change the machine state)

In-order issue with in-order completion

In-order issue with out-of-order completion

Out-of-order issue with out-of-order completion and in-order commit

Out-of-order issue with out-of-order completion


In-Order Issue with In-Order Completion

- Simplest policy is to issue instructions in exact program order and to complete them in the same order they were fetched (i.e., in program order)

In-Order Issue with In-Order Completion (Ex.)


- Assume a pipelined processor that can fetch and decode two instructions per cycle, that has three functional units (a single cycle adder, a single cycle shifter, and a two cycle multiplier), and that can complete (and write back) two results per cycle
- Instruction sequence:
 - l₁: needs two execute cycles (a multiply)
 - l₂
 - l₃
 - l₄: needs the same function unit as l₃
 - l₅: needs data value produced by l₄
 - l₆: needs the same function unit as l₅

In-Order Issue, In-Order Completion Example




Pentium Retrospective

- Limited in performance by “front end”
 - Has to support variable-length instrs and segments
- Supporting all x86 features tough!
 - 30% of transistors are for legacy support
 - Up to 40% in Pentium Pro!
 - Down to 10% in P4
 - Microcode ROM is huge



Pentium Retrospective


- Pentium is in-order issue, in-order complete
- “Static scheduling” by the dispatch logic:
 - Fetch/dispatch/execute/retire: all in order
- Drawbacks:
 - Adapts poorly to dynamic code stream
 - Adapts poorly to future hardware
 - What if we had 3 pipes not 2?



In-Order Issue with Out-of-Order Completion

- With out-of-order completion, a later instruction may complete **before** a previous instruction
- Out-of-order completion is used in single-issue pipelined processors to improve the performance of long-latency operations such as divide
- When using out-of-order completion instruction issue is **stalled** when there is a resource conflict (e.g., for a functional unit) or when the instructions ready to issue need a result that has not yet been computed

I0I-00C Example



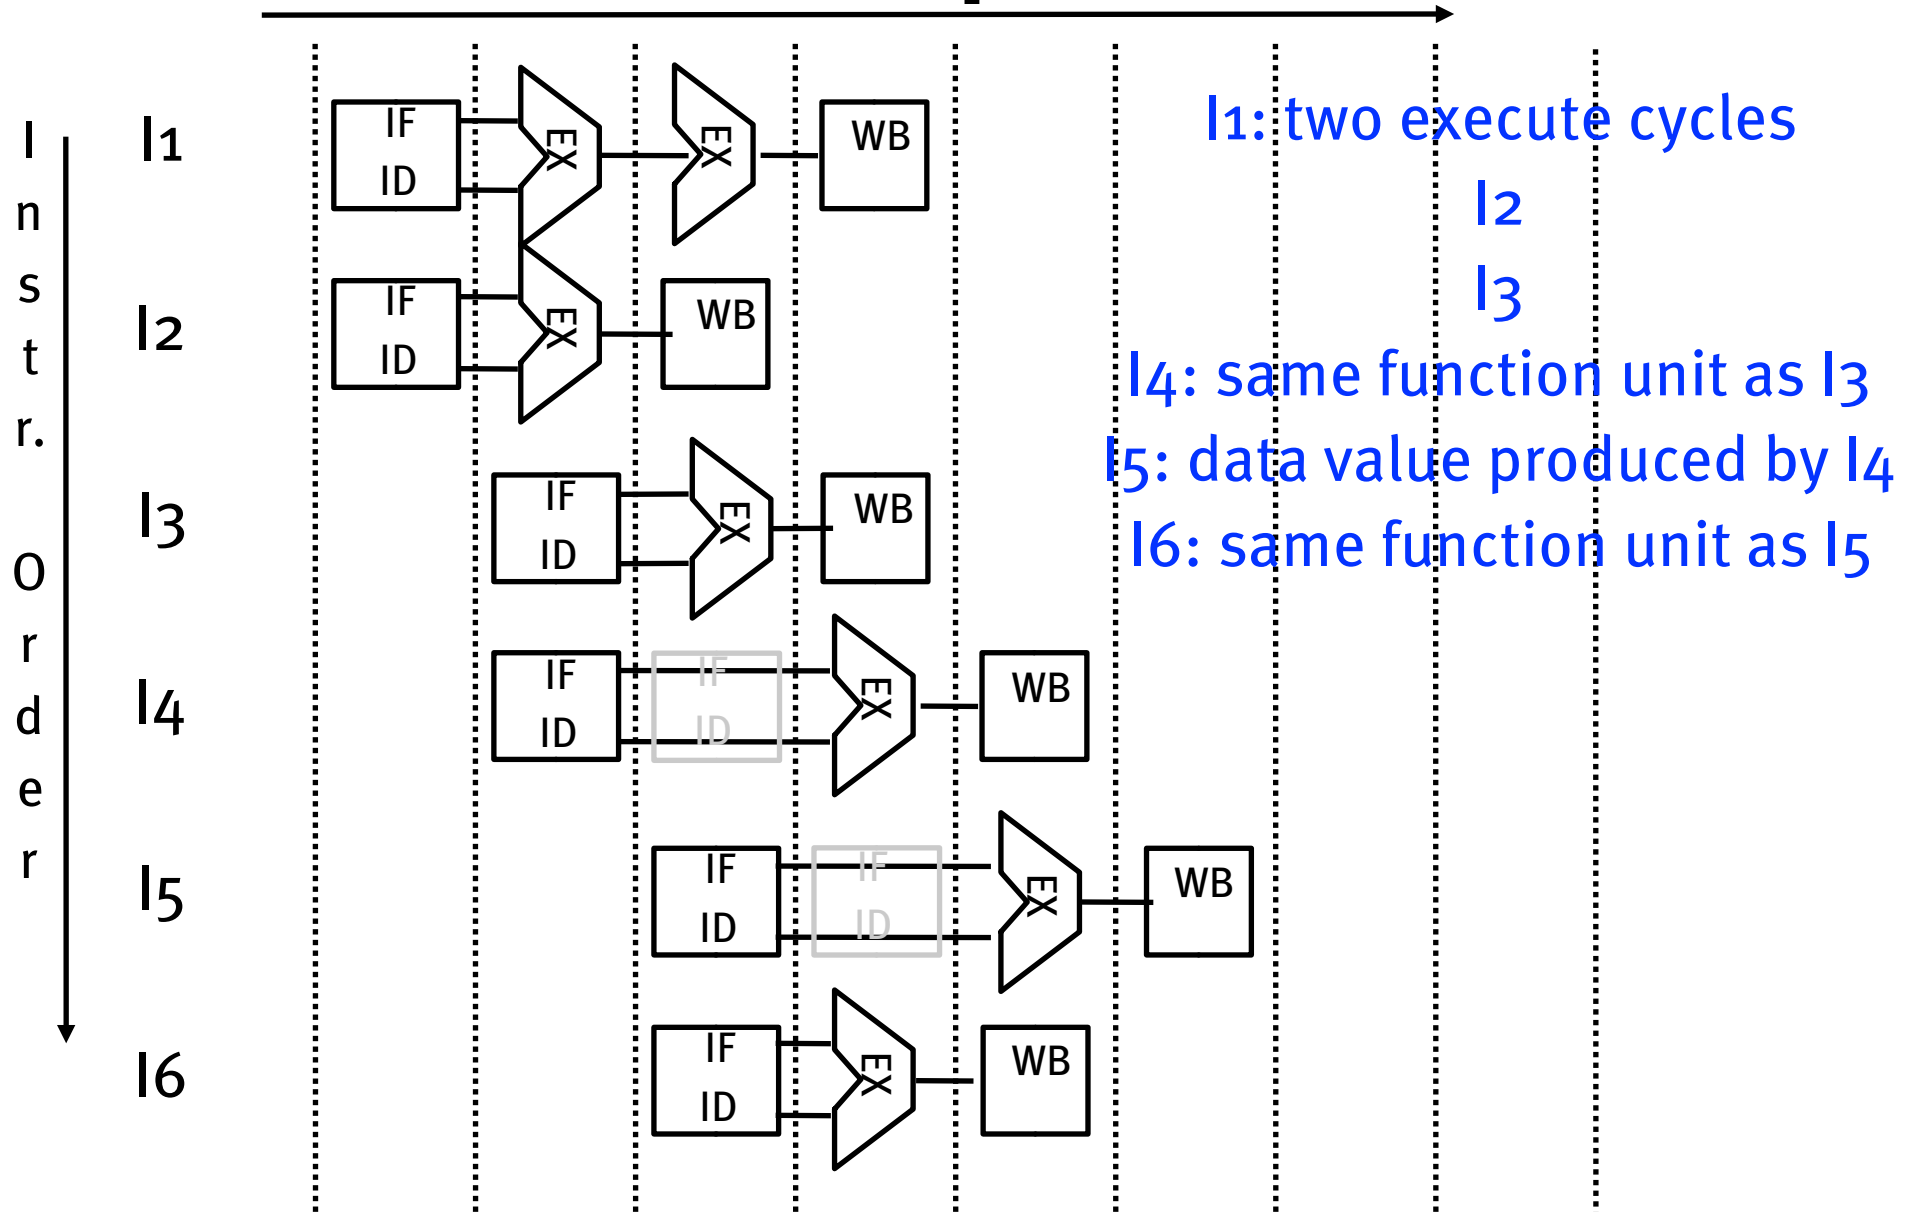
Handling Output Dependencies

- There is one more situation that stalls instruction issuing with IOI-OOC, assume
 - I₁ – writes to R₃
I₂ – writes to R₃
I₅ – reads R₃
- If the I₁ write occurs after the I₂ write, then I₅ reads an incorrect value for R₃
- I₂ has an output dependency on I₁—write before write (we would say “write after write conflict”)
- The issuing of I₂ would have to be stalled if its result might later be overwritten by an previous instruction (i.e., I₁) that takes longer to complete—the stall happens before instruction issue
- While IOI-OOC yields higher performance, it requires more dependency checking hardware (both read-before-write and write-before-write)

Out-of-Order Issue with Out-of-Order Completion

- With in-order issue the processor stops decoding instructions whenever a decoded instruction has a resource conflict or a data dependency on an issued, but uncompleted instruction
 - The processor is not able to look beyond the conflicted instruction even though more downstream instructions might have no conflicts and thus be issueable
- Fetch and decode instructions beyond the conflicted one (“**instruction window**”: Tetris), store them in an instruction buffer (as long as there's room), and flag those instructions in the buffer that don't have resource conflicts or data dependencies
- Flagged instructions are then issued from the buffer without regard to their program order

001-00C Example



Dependency Examples

- $R_3 := R_3 * R_5$
 $R_4 := R_3 + 1$
 $R_3 := R_5 + 1$

True data dependency (RAW)

Output dependency (WAW)

Antidependency (WAR)

Antidependencies

- With OOI also have to deal with data antidependencies – when a later instruction (that completes earlier) produces a data value that destroys a data value used as a source in an earlier instruction (that issues later)
- The constraint is similar to that of true data dependencies, except reversed
- Instead of the later instruction using a value (not yet) produced by an earlier instruction (read before write), the later instruction produces a value that destroys a value that the earlier instruction (has not yet) used (write before read)

Dependencies Review

- Each of the three data dependencies ...
 - True data dependencies (read before write)
 - Antidependencies (write before read)
 - Output dependencies (write before write)
- } storage conflicts
- ... manifests itself through the use of registers (or other storage locations)
 - True dependencies represent the flow of data and information through a program
 - Anti- and output dependencies arise because the limited number of registers mean that programmers reuse registers for different computations
 - When instructions are issued out-of-order, the correspondence between registers and values breaks down and the values conflict for registers

Conflict example

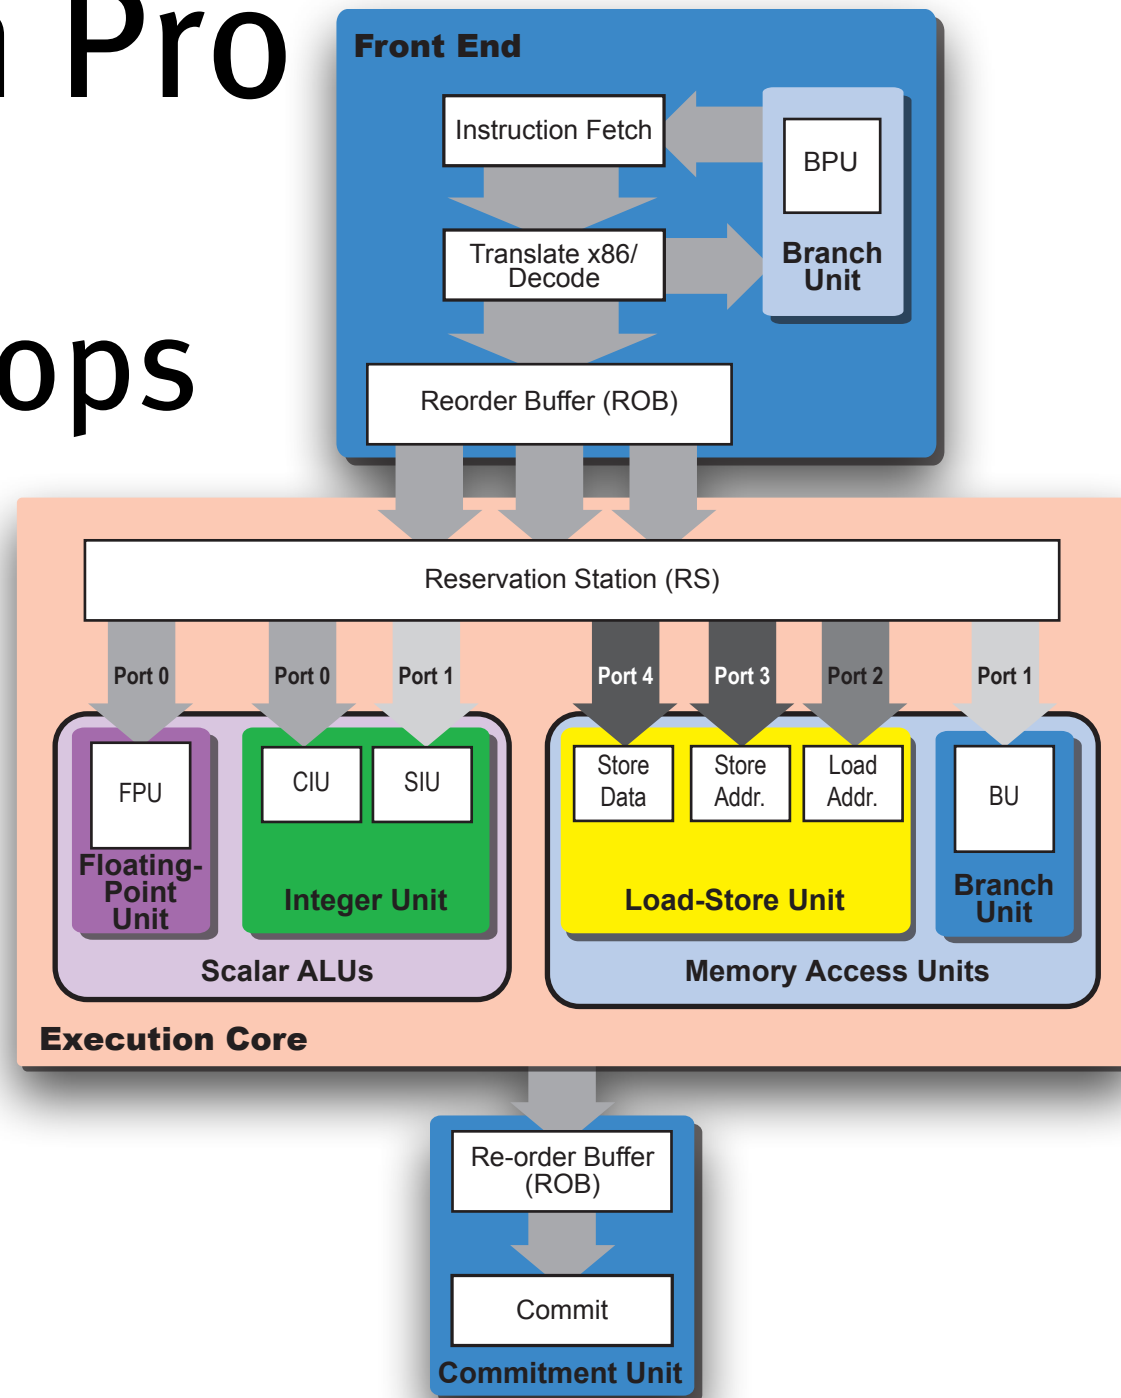
- (1) $R_3 := R_3 * R_5$
(2) $R_4 := R_3 + 1$
(3) $R_3 := R_5 + 1$
 - (3) must ensure that (1) and (2) read R_3 before (3) writes it
 - But the value in (3)'s R_3 has nothing to do with the value in (1) and (2)'s R_3 ! Shouldn't we be able to issue that instruction?

Storage Conflicts and Register Renaming

- Storage conflicts can be reduced (or eliminated) by increasing or duplicating the troublesome resource
- Provide additional registers that are used to reestablish the correspondence between registers and values
 - Allocated dynamically by the hardware in SS processors
- Register renaming — the processor renames the original register identifier in the instruction to a new register (one not in the visible register set)
- | | |
|--------------------|-----------------------------|
| $R_3 := R_3 * R_5$ | $R_{3b} := R_{3a} * R_{5a}$ |
| $R_4 := R_3 + 1$ | $R_{4a} := R_{3b} + 1$ |
| $R_3 := R_5 + 1$ | $R_{3c} := R_{5a} + 1$ |
- The hardware that does renaming assigns a “replacement” register from a pool of free registers and releases it back to the pool when its value is superseded and there are no outstanding references to it


Pentium Pro

uops



Pentium Pro

- | | |
|-----------------------|--------------|
| 1. Fetch | In order |
| 2. Decode/dispatch | In order |
| 3. Issue | Reorder |
| 4. Execute | Out of order |
| 5. Complete | Reorder |
| 6. Writeback (commit) | In order |

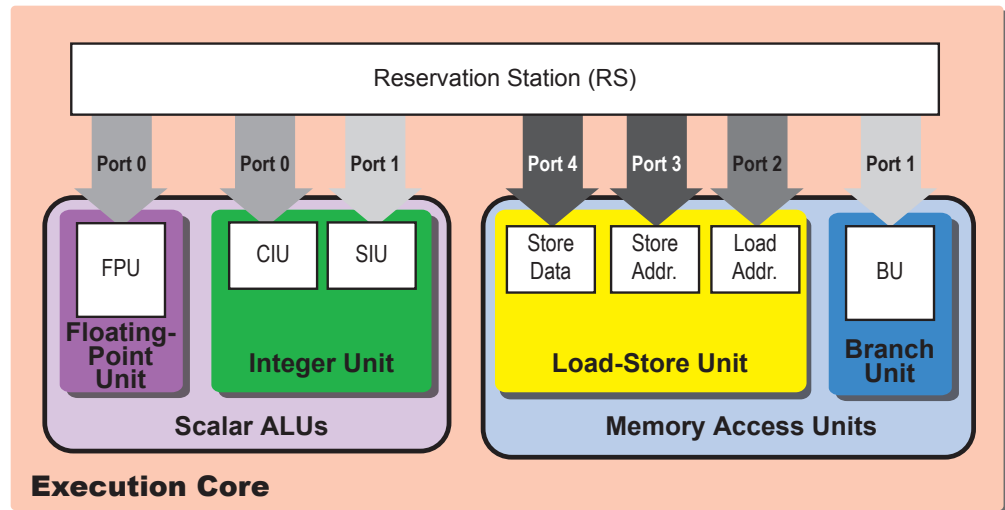


P6 Pipeline

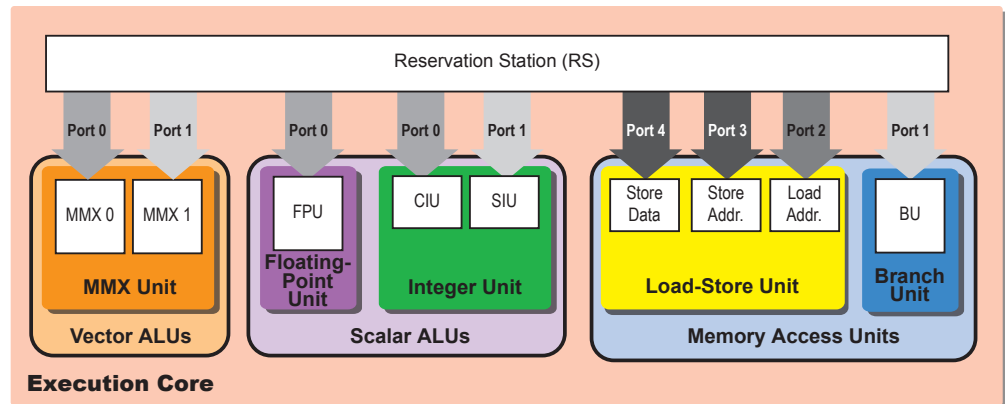
- Instruction fetch, BTB access (3.5 stages)
 - 2 cycles for instruction fetch
- Decode, x86- \rightarrow uops (2.5 stages)
- Register rename (1 stage)
- Write to reservation station (1 stage)
- Read from reservation station (1 stage)
- Execute (1+ stages)
- Commit (2 stages)

Pentium Pro backends


- Pentium Pro



- Pentium 2




- Pentium 3



Apple “Cyclone” A7 (2014)

Apple Cyclone



Apple Custom CPU Core Comparison


	Apple A6	Apple A7
CPU Codename	Swift	Cyclone
ARM ISA	ARMv7-A (32-bit)	ARMv8-A (32/64-bit)
Issue Width	3 micro-ops	6 micro-ops
Reorder Buffer Size	45 micro-ops	192 micro-ops
Branch Mispredict Penalty	14 cycles	16 cycles (14 - 19)
Integer ALUs	2	4
Load/Store Units	1	2
Load Latency	3 cycles	4 cycles
Branch Units	1	2
Indirect Branch Units	0	1
FP/NEON ALUs	?	3
L1 Cache	32KB I\$ + 32KB D\$	64KB IS + 64KB D\$
L2 Cache	1MB	1MB
L3 Cache	-	4MB

AMD Zen (2016)



ZEN MICROARCHITECTURE

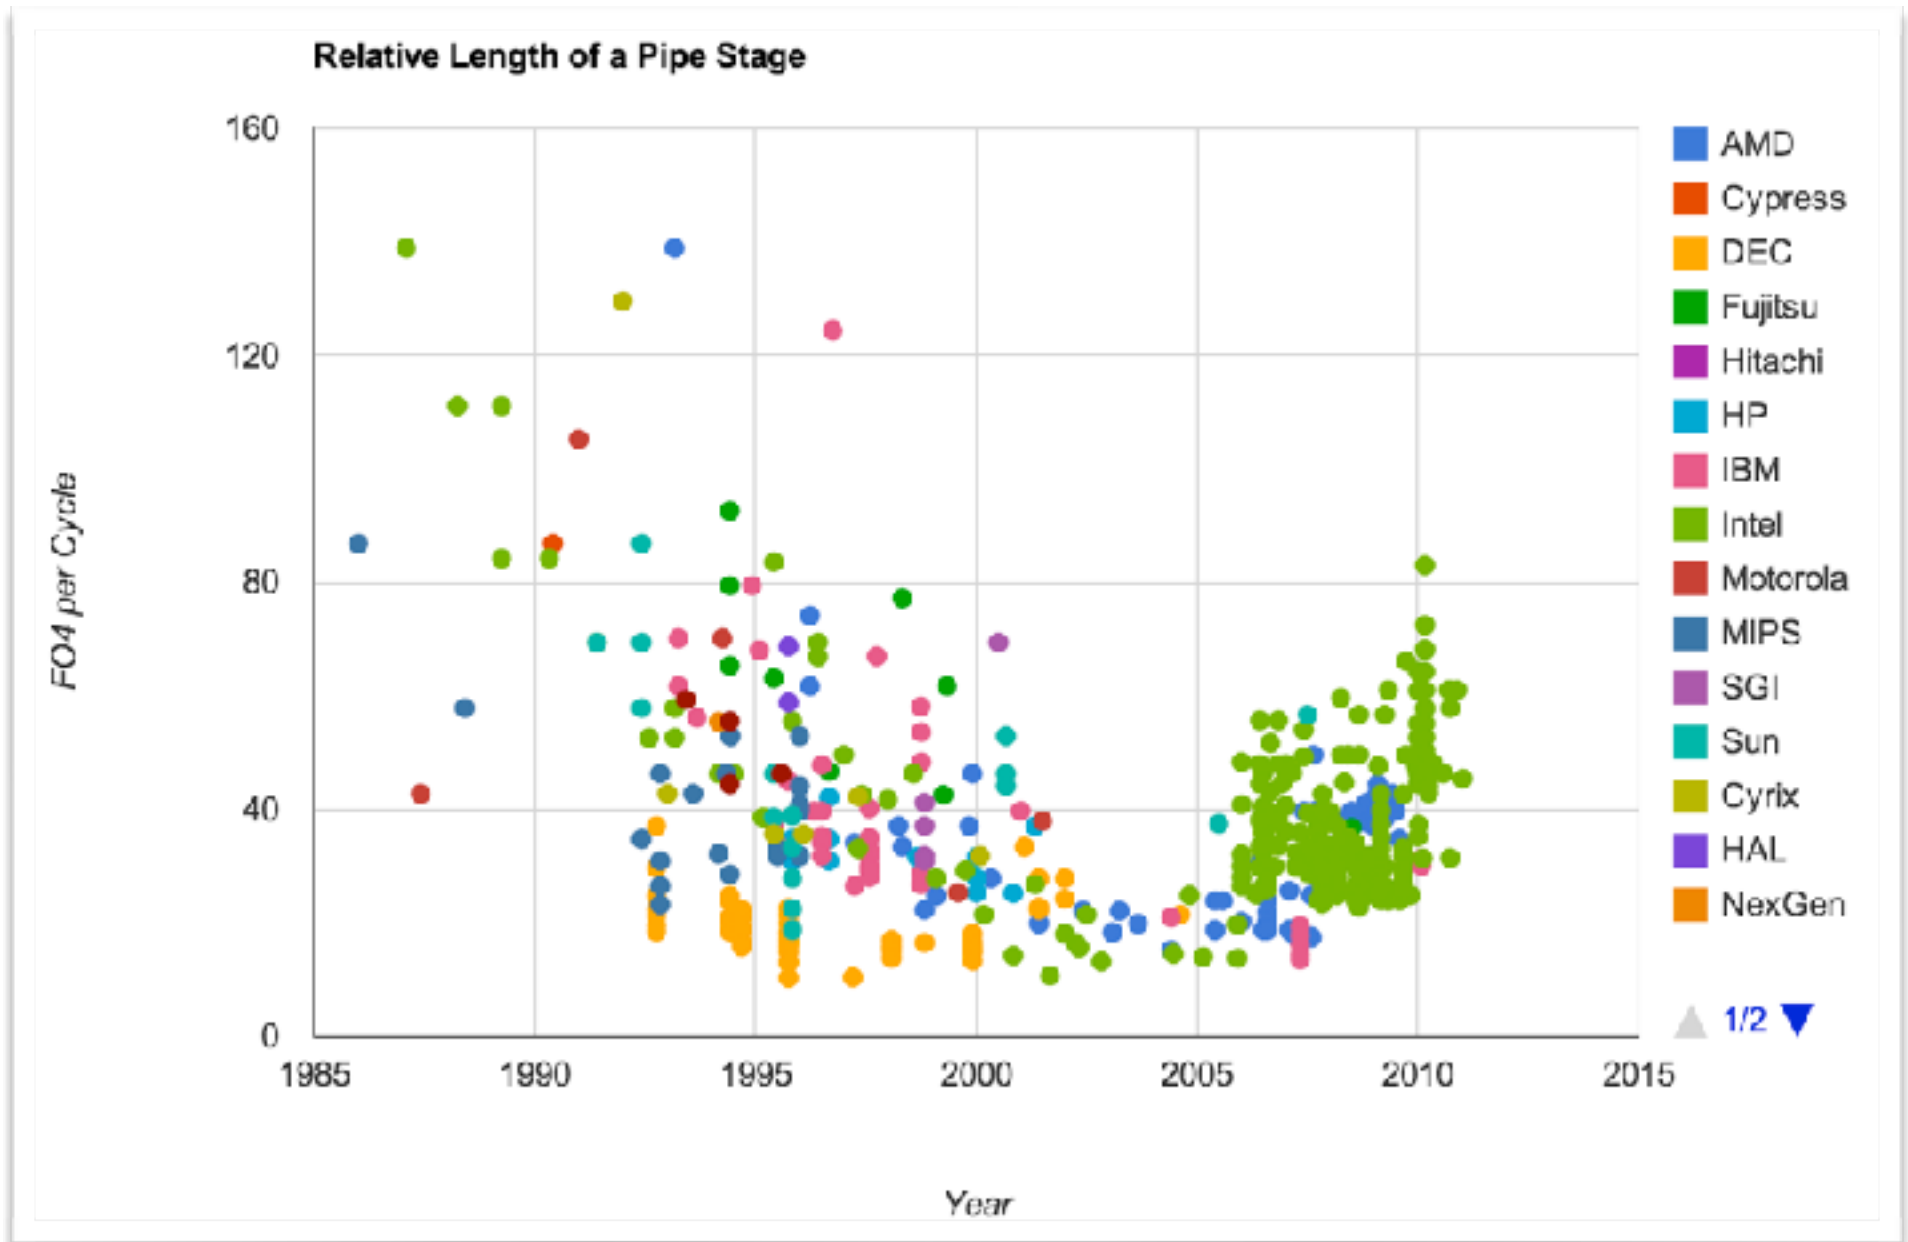
- Fetch Four x86 instructions
- Op Cache instructions
- 4 Integer units
 - Large rename space – 168 Registers
 - 192 instructions in flight/8 wide retire
- 2 Load/Store units
 - 72 Out-of-Order Loads supported
- 2 Floating Point units x 128 FMACs
 - built as 4 pipes, 2 Fadd, 2 Fmul
- I-Cache 64K, 4-way
- D-Cache 32K, 8-way
- L2 Cache 512K, 8-way
- Large shared L3 cache
- 2 threads per core



Where Do We Get ILP?

- All of these techniques require that we have ample instruction level parallelism
 - Original P4 has 20 stages, 6 μ ops per cycle
 - Lots of instructions in flight!

Hardware limits to superpipelining?



http://cpudb.stanford.edu/visualize/fo4_per_cycle

VLIW Beginnings

- VLIW: Very Long Instruction Word

[4] J. A. Fisher, "Very long instruction word architectures and the ELI-512," in *Proc. 10th Symp. Comput. Architecture*, IEEE, June 1983, pp. 140–150.

- Josh Fisher: idea grew out of his Ph.D (1979) in compilers
- Led to a startup (MultiFlow) whose computers worked, but which went out of business ... the ideas remain influential.

History of VLIW Processors

- Started with (horizontal) microprogramming
 - Very wide microinstructions used to directly generate control signals in single-issue processors (e.g., IBM 360 series)
- VLIW for multi-issue processors first appeared in the Multiflow and Cydrome (in the early 1980's)
- Current commercial VLIW processors
 - Intel i860 RISC (dual mode: scalar and VLIW)
 - Intel I-64 ([EPIC](#): Itanium and Itanium 2) [future lecture]
 - Transmeta Crusoe [future case study]
 - Lucent/Motorola StarCore, ADI TigerSHARC, Infineon (Siemens) Carmel

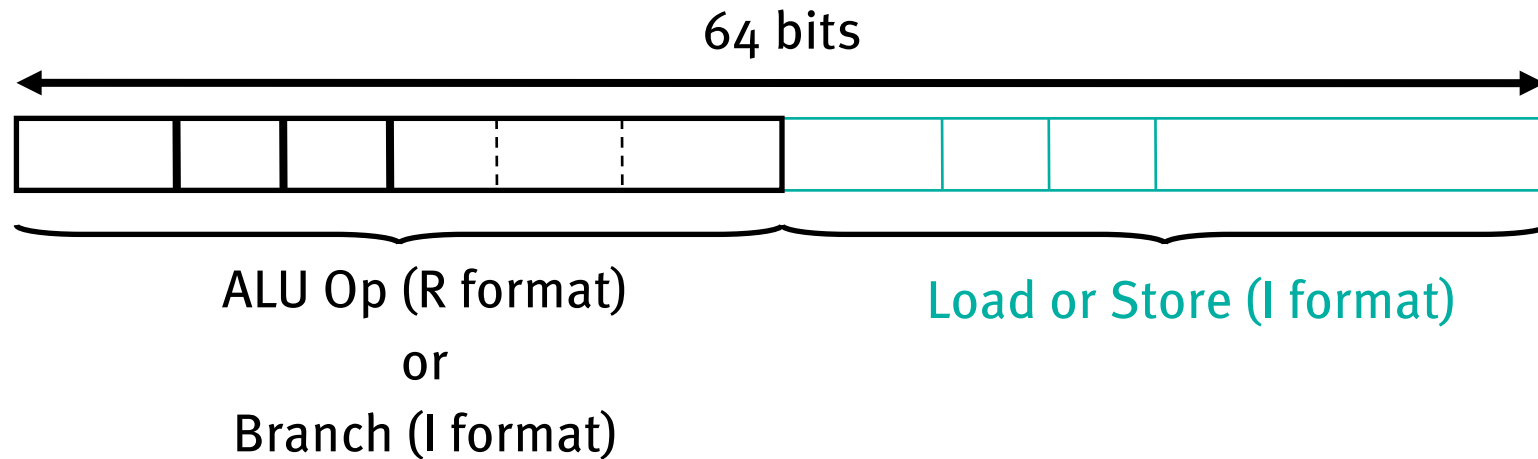
Static Multiple Issue Machines (VLIW)

- Static multiple-issue processors (aka VLIW) use the compiler to decide which instructions to issue and execute simultaneously
- Issue packet—the set of instructions that are bundled together and issued in one clock cycle—think of it as one large instruction with multiple operations
- The mix of instructions in the packet (bundle) is usually restricted—a single “instruction” with several predefined fields
- The compiler does static branch prediction and code scheduling to reduce (ctrl) or eliminate (data) hazards

Static Multiple Issue Machines (VLIW)

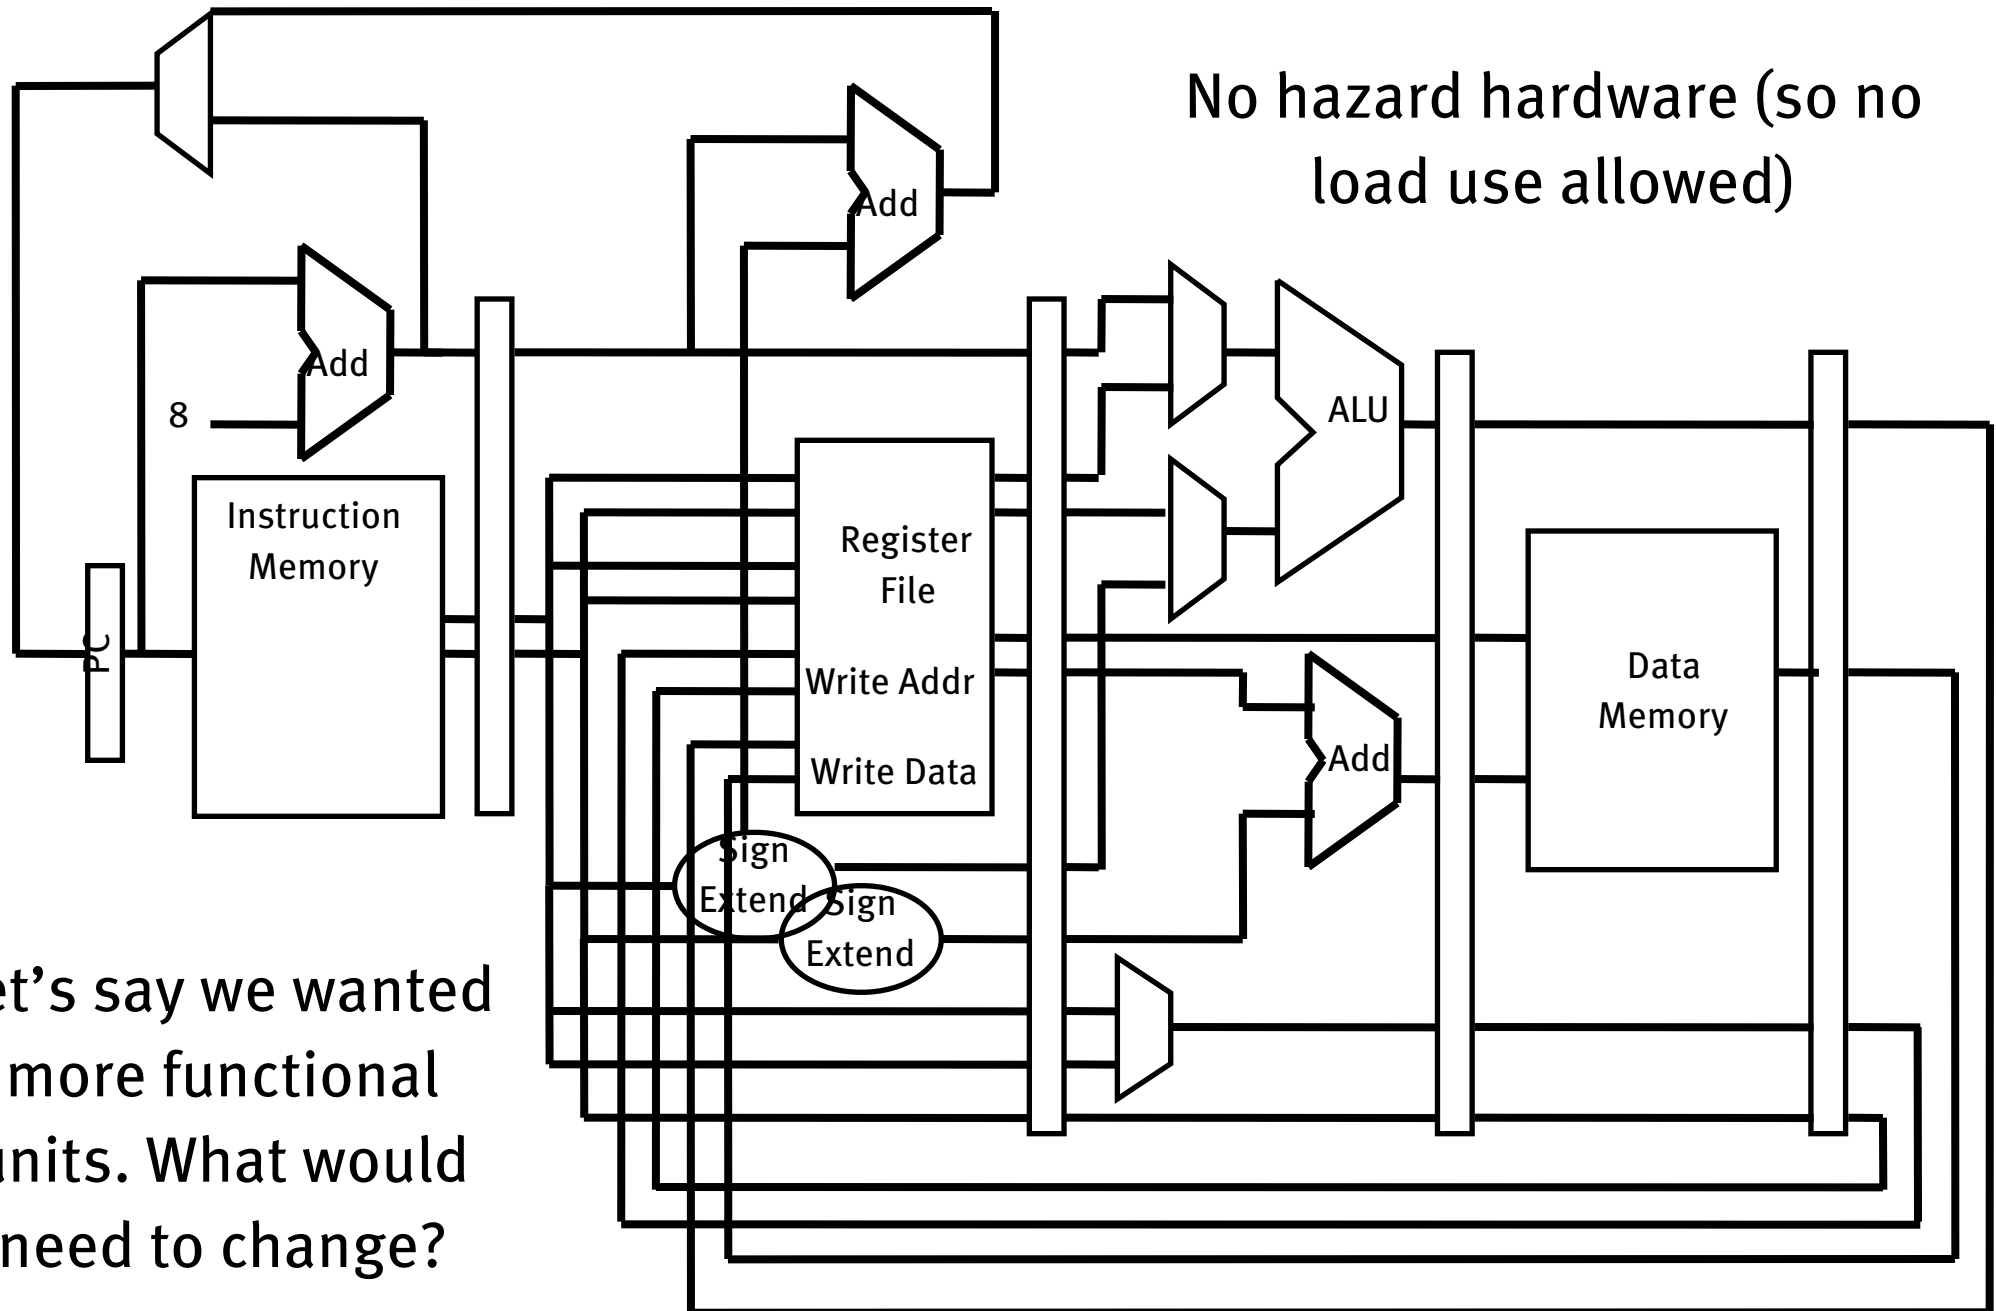
- VLIWs have
 - Multiple functional units (like SS processors)
 - Multi-ported register files (again like SS processors)
 - Wide program bus

An Example: A VLIW MIPS



- Consider a 2-issue MIPS with a 2 instr bundle
- Instructions are always fetched, decoded, and issued in pairs
 - What happens if one instr isn't used?
 - What does the register file have to support?
 - What other hw must we add?

A MIPS VLIW (2-issue) Datapath



Code Scheduling Example

- Consider the following loop code:

```
lp:  lw      $to,o($s1)    # $to=array element
     addu    $to,$to,$s2    # add scalar in $s2
     sw      $to,o($s1)    # store result
     addi    $s1,$s1,-4    # decrement pointer
     bne     $s1,$o,lp     # branch if $s1 != 0
```

```
i = n;
do {
    a[i] += const;
} while (--i != 0);
```

- Must “schedule” the instructions to avoid pipeline stalls
 - All instructions can forward results to be available the next cycle EXCEPT must separate load use instructions from their loads by one cycle
 - Instructions in one bundle must be independent
 - Notice that the first two instructions have a load use dependency, the next two and last two have data dependencies
 - Assume branches are perfectly predicted by the hardware

The Scheduled Code (Not Unrolled)

	ALU or branch	Data transfer	CC
lp:			1
			2
			3
			4
			5

- How many clock cycles?
- How many instructions?
- CPI? Best case?
- IPC? Best case?

Loop Unrolling

- Loop unrolling—multiple copies of the loop body are made and instructions from different iterations are scheduled together as a way to increase ILP
- Apply loop unrolling (4 times for our example) and then schedule the resulting code
 - Eliminate unnecessary loop overhead instructions
 - Schedule so as to avoid load use hazards
- During unrolling the compiler applies register renaming to eliminate all data dependencies that are not true dependencies

Unrolled Code Example

- lp: lw \$t0,0(\$s1) # \$t0=array element
lw \$t1,-4(\$s1) # \$t1=array element
lw \$t2,-8(\$s1) # \$t2=array element
lw \$t3,-12(\$s1) # \$t3=array element
addu \$t0,\$t0,\$s2 # add scalar in \$s2
addu \$t1,\$t1,\$s2 # add scalar in \$s2
addu \$t2,\$t2,\$s2 # add scalar in \$s2
addu \$t3,\$t3,\$s2 # add scalar in \$s2
sw \$t0,0(\$s1) # store result
sw \$t1,-4(\$s1) # store result
sw \$t2,-8(\$s1) # store result
sw \$t3,-12(\$s1) # store result
addi \$s1,\$s1,-16 # decrement pointer
bne \$s1,\$0,lp # branch if \$s1 != 0

The Scheduled Code (Unrolled)

	ALU or branch	Data transfer	CC
lp:	addi \$s1,\$s1,-16	lw \$t0,0(\$s1)	1
		lw \$t1,12(\$s1)	2
	addu \$t0,\$t0,\$s2	lw \$t2,8(\$s1)	3
	addu \$t1,\$t1,\$s2	lw \$t3,4(\$s1)	4
	addu \$t2,\$t2,\$s2	sw \$t0,16(\$s1)	5
	addu \$t3,\$t3,\$s2	sw \$t1,12(\$s1)	6
		sw \$t2,8(\$s1)	7
	bne \$s1,\$0,lp	sw \$t3,4(\$s1)	8

- Eight clock cycles to execute 14 instructions for a
 - CPI of 0.57 (versus the best case of 0.5)
 - IPC of 1.8 (versus the best case of 2.0)

What does N = 14 assembly look like?

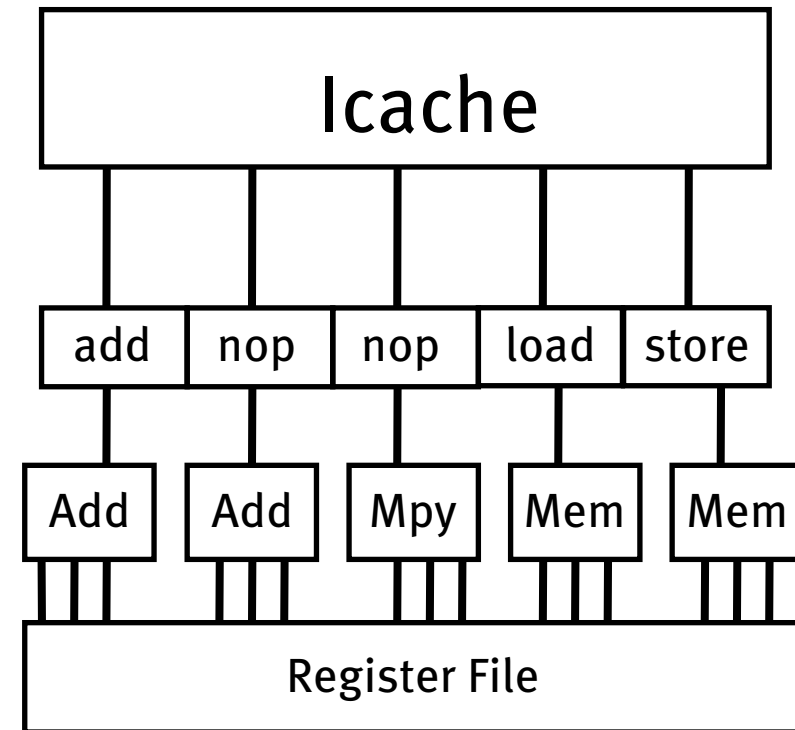
- Two instructions from a scientific benchmark (Linpack) for a MultiFlow CPU with 14 operations per instruction.

instr	cl0	ialu0e	st.64	sb1.r0,r2,17#144
	cl0	ialu1e	cgt.s32	li1bb.r4,r34,6#31
	cl0	falu0e	add.f64	lsb.r4,r8,r0
	cl0	falu1e	add.f64	lsb.r6,r40,r32
	cl0	ialu0l	dld.64	fb1.r4,r2,17#208
	cl1	ialu0e	dld.64	fb1.r34,r1,17#216
	cl1	ialu1e	cgt.s32	li1bb.r3,r32,zero
	cl1	falu0e	add.f64	lsb.r4,r8,r6
	cl1	falu1e	add.f64	lsb.r6,r40,r38
	cl1	ialu0l	st.64	sb1.r2,r1,17#152
	cl1	ialu1l	add.u32	lib.r32,r36,6#32
	cl1	br	true and r3	L23?3
	cl0	br	false or r4	L24?3;

instr	cl0	ialu0e	dld.64	fb0.r0,r2,17#224
	cl0	ialu1e	cgt.s32	li1bb.r3,r34,6#30
	cl0	falu0e	mpy.f64	lfb.r10,r2,r10
	cl0	falu1e	mpy.f64	lfb.r42,r34,r42
	cl0	ialu0l	st.64	sb0.r4,r2,17#160
	cl1	ialu0e	dld.64	fb0.r32,r1,17#232
	cl1	ialu1e	cgt.s32	li1bb.r4,r35,6#29
	cl1	falu0e	mpy.f64	lfb.r10,r0,r10
	cl1	falu1e	mpy.f64	lfb.r42,r32,r42
	cl1	ialu0l	st.64	sb0.r6,r1,17#168
	cl1	ialu1l	bor.32	ib0.r32,zero,r32
	cl1	br	false or r4	L25?3
	cl0	br	true and r3	L26?3;

Defining Attributes of VLIW

- Compiler:
 - 1. MultiOp: instruction containing multiple independent operations
 - 2. Specified number of resources of specified types
 - 3. Exposed, architectural latencies



VLIW instruction =
5 independent
operations

Compiler Support for VLIW Processors

- The compiler packs groups of **independent** instructions into the bundle
 - Because branch prediction is not perfect, done by code re-ordering (trace scheduling)
 - We'll cover this in a future lecture
- The compiler uses loop unrolling to expose more ILP
- The compiler uses register renaming to solve name dependencies and ensures no load use hazards occur

Compiler Support for VLIW Processors

- While superscalars use dynamic prediction, VLIWs primarily depend on the compiler for extracting ILP
- Loop unrolling reduces the number of conditional branches
- Predication eliminates if-the-else branch structures by replacing them with predicated instructions
 - We'll cover this in a future lecture as well
- The compiler predicts memory bank references to help minimize memory bank conflicts

VLIW Advantages

- Advantages
 - Simpler hardware (potentially less power hungry)
 - Potentially more scalable
 - Allow more instr's per VLIW bundle and add more FUs

VLIW Disadvantages

- Programmer/compiler complexity and longer compilation times
 - Deep pipelines and long latencies can be confusing (making peak performance elusive)
- Lock step operation, i.e., on hazard all future issues stall until hazard is resolved (hence need for predication)
- Object (binary) code incompatibility
- Needs lots of program memory bandwidth
- Code bloat
 - Noops are a waste of program memory space
 - Loop unrolling to expose more ILP uses more program memory space

Review: Multi-Issue Datapath Responsibilities

- Must handle, with a combination of hardware and software
 - Data dependencies – aka data hazards
 - True data dependencies (read after write)
 - Use data forwarding hardware
 - Use compiler scheduling
 - Storage dependence (aka name dependence)
 - Use register renaming to solve both
 - Antidependencies (write after read)
 - Output dependencies (write after write)

Review: Multi-Issue Datapath Responsibilities

- Must handle, with a combination of hardware and software
- Procedural dependencies—aka control hazards
 - Use aggressive branch prediction (speculation)
 - Use predication
 - Future lecture

Review: Multi-Issue Datapath Responsibilities

- Must handle, with a combination of hardware and software
- Resource conflicts—aka structural hazards
 - Use resource duplication or resource pipelining to reduce (or eliminate) resource conflicts
 - Use arbitration for result and commit buses and register file read and write ports

Review: Multiple-Issue Processor Styles

- Dynamic multiple-issue processors (aka **superscalar**)
 - Decisions on which instructions to execute simultaneously (in the range of 2 to 8 in 2005) are being made dynamically (at run time by the **hardware**)
 - E.g., IBM Power 2, Pentium 4, MIPS R10K, HP PA 8500 IBM
- Static multiple-issue processors (aka **VLIW**)
 - Decisions on which instructions to execute simultaneously are being made statically (at compile time by the **compiler**)
 - E.g., Intel Itanium and Itanium 2 for the IA-64 ISA – EPIC (Explicit Parallel Instruction Computer)
 - 128 bit “bundles” containing 3 instructions each 41 bits + 5 bit template field (specifies which FU each instr needs)
 - Five functional units (IntALU, MMedia, DMem, FPALU, Branch)
 - Extensive support for speculation and predication

CISC vs RISC vs SS vs VLIW

	CISC	RISC	Super-scalar	VLIW
Instr size	variable size	fixed size	fixed size	fixed size (but large)
Instr format	variable format	fixed format	fixed format	fixed format
Registers	few, some special	many GP	GP and rename	many, many GP
Memory reference	embedded in many instrs	load/store	load/store	load/store
Key Issues	decode complexity	data forwarding, hazards	hardware dependency resolution	(compiler) code scheduling
Instruction flow		