

Lecture 19

Wrapup

EEC 171 Parallel Architectures

John Owens

UC Davis

Credits

- © John Owens / UC Davis 2007–17.
- Thanks to many sources for slide material: Computer Organization and Design (Patterson & Hennessy) © 2005, Computer Architecture (Hennessy & Patterson) © 2007, Inside the Machine (Jon Stokes) © 2007, © Dan Connors / University of Colorado 2007, © Kathy Yelick / UCB 2007, © Wen-Mei Hwu/David Kirk, University of Illinois 2007, © David Patterson / UCB 2003–7, © John Lazzaro / UCB 2006, © Mary Jane Irwin / Penn State 2005, © John Kubiawicz / UCB 2002, © Krste Asinovic/Arvind / MIT 2002, © Morgan Kaufmann Publishers 1998.


Also thanks to ...

- Kathy Yelick, for her slides on “A Berkeley View on the Parallel Computing Landscape”
- Horst Simon, for his slides on “Supercomputers and Superintelligence”
- A prepublication copy of “The Future of Computing Performance” (2010) by the National Research Council that I found online

New: Power Wall


Can put more transistors on a chip than can afford to turn on

Scaling clock speed (business as usual) will not work




Very Old: Multiplies Slow, Loads fast


- Design algorithms to reduce floating point operations
- Machines measured on peak flop/s



Performance Development




Average System Age



New: Memory Performance is Key

Ever-growing processor-memory performance gap




- Total chip performance still growing with Moore's Law
- Bandwidth rather than latency will be growing concern

New: Clock Scaling Bonanza Has Ended


- Chip density is continuing increase ~2x every 2 years
 - Clock speed is not
 - Number of processor cores may double instead
- There is little or no hidden parallelism (ILP) to be found
- Parallelism must be exposed to and managed by software

Source: Intel, Microsoft (Sutter) and Stanford (Olukotun, Hammond)

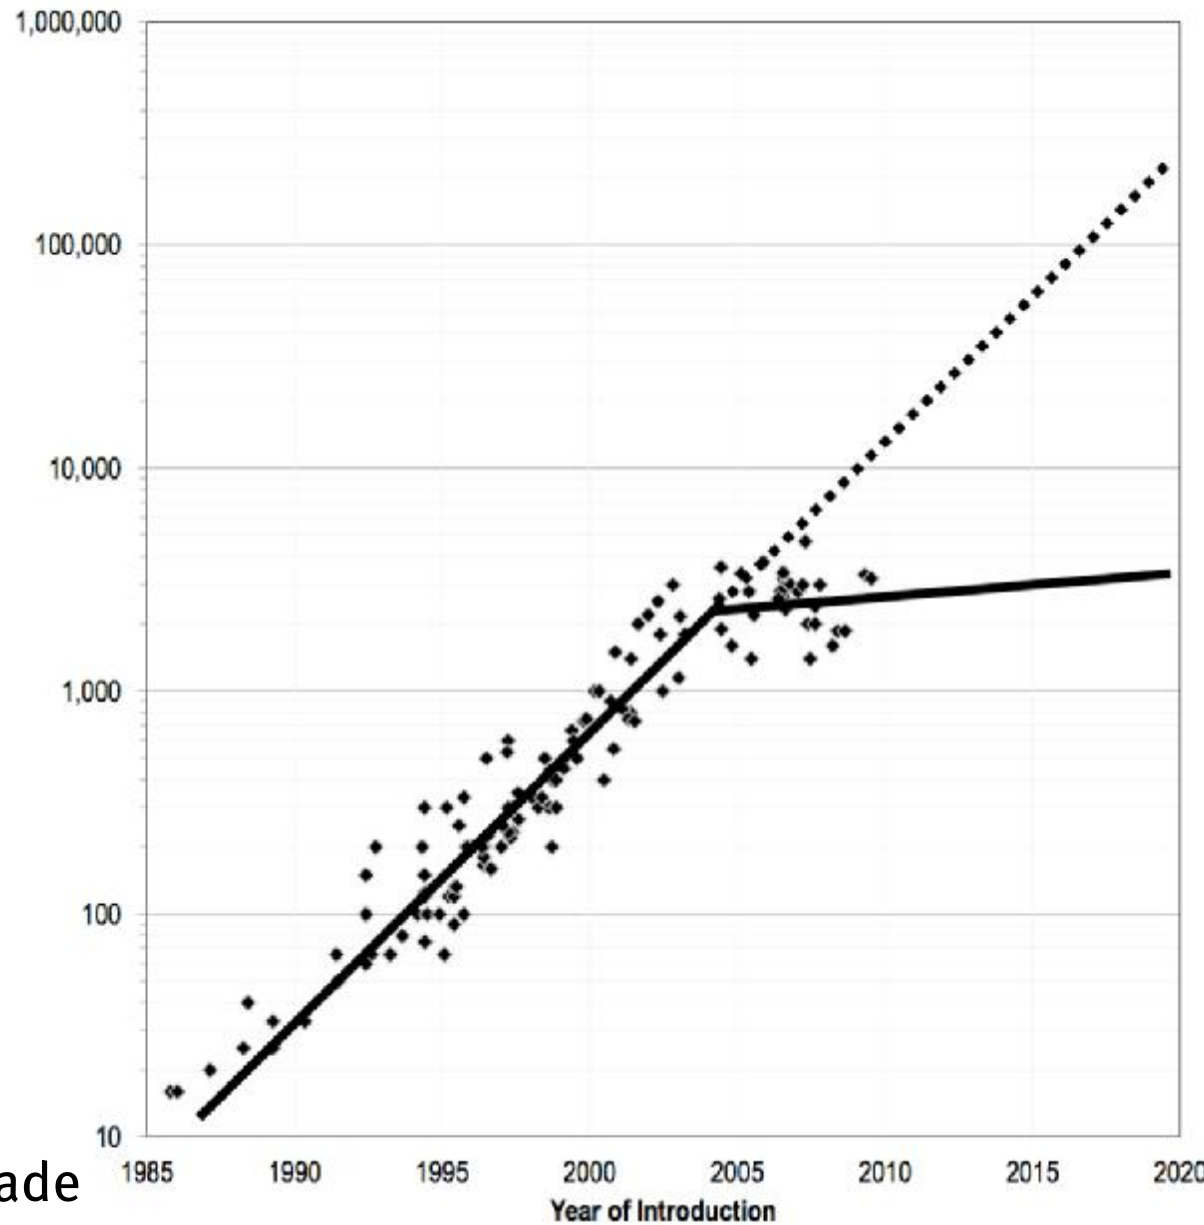


Single-core performance slowing

SPEC Integer Performance (single proc. x86)



Jim Larus, Microsoft, from a talk at UC Davis in May 2009



NRC report, SPECint2000

From 100x/decade to 2x/decade

Old: Parallelism only for High End Computing

Address http://www.parelogos.com/DeadS_per/1-index.html

The Dead Supercomputer Society



The Passing of a Golden Age?

From the construction of the first programmed computers until the mid 1990s, there was always room in the computer industry for someone with a clever, if sometimes challenging, idea on how to make a more powerful machine. Computing became strategic during the Second World War, and remained so during the Cold War that followed. High-performance computing is essential to any modern nuclear weapons program, and a computer technology "race" was a logical corollary to the arms race. While powerful computers are of great value to a number

New: Parallelism by Necessity

“This shift toward increasing parallelism is not a triumphant stride forward based on breakthroughs in novel software and architectures for parallelism; instead, this **plunge into parallelism is actually a retreat from even greater challenges that thwart efficient silicon implementation of traditional uniprocessor architectures.**”


Kurt Keutzer, Berkeley View, December 2006

- HW/SW Industry bet its future that breakthroughs will appear before it's too late
 - David Patterson: “Industry has already thrown the Hail Mary pass ... but nobody is running yet.”

Conventional Wisdom (CW) in Computer Architecture

1. *Old CW*: Power is free, but transistors expensive
 - *New CW*: Power wall Power expensive, transistors “free”
 - Can put more transistors on a chip than have the power to turn on
2. *Very Old CW*: Multiplies slow, but loads fast
 - *New CW*: Memory wall Loads slow, multiplies fast
 - 200 clocks to DRAM, but even FP multiplies only 4 clocks
3. *Old CW*: More ILP via compiler/architecture innovation
 - Branch prediction, speculation, Out-of-order execution, VLIW, ...
 - *New CW*: ILP wall Diminishing returns on more ILP
4. *Old CW*: 2X CPU Performance every 18 months
 - *New CW*: *Power + Memory + ILP Walls = Brick Wall*
5. *Old CW*: Parallelism is only for Scientific fringe
 - *New CW*: *Parallelism is everywhere*

7 Questions for Parallelism



■ Applications:

1. What are the apps?
2. What are kernels of apps?

■ Hardware:


3. What are the HW building blocks?
4. How to connect them?

■ Programming Model & Systems Software:

5. How to describe apps and kernels?
6. How to program the HW?

■ Evaluation:

7. How to measure success?



(Inspired by a view of the Golden Gate Bridge from Berkeley)

What do you see as future
directions for higher-performance
computing?


Example Applications in Health

- Imagine a “digital body double”
 - 3D image-based medical record
 - Includes diagnostic, pathologic, and other information
- Used for:
 - Diagnosis
 - Less invasive surgery-by-robot
 - Experimental treatments
 - Real-time diet and exercise recommendations




Existing simulations

- Heart
- Lung
- Brain
- Kidney
- Bone mass



RMS Applications



Recognition



What is a hedge?
What is the interest rate?

Mining



Is there a hedging
opportunity here?

Synthesis




What if interest rates
were to go up?

Recognition




What is a tumor?

Mining



Is there a tumor here?

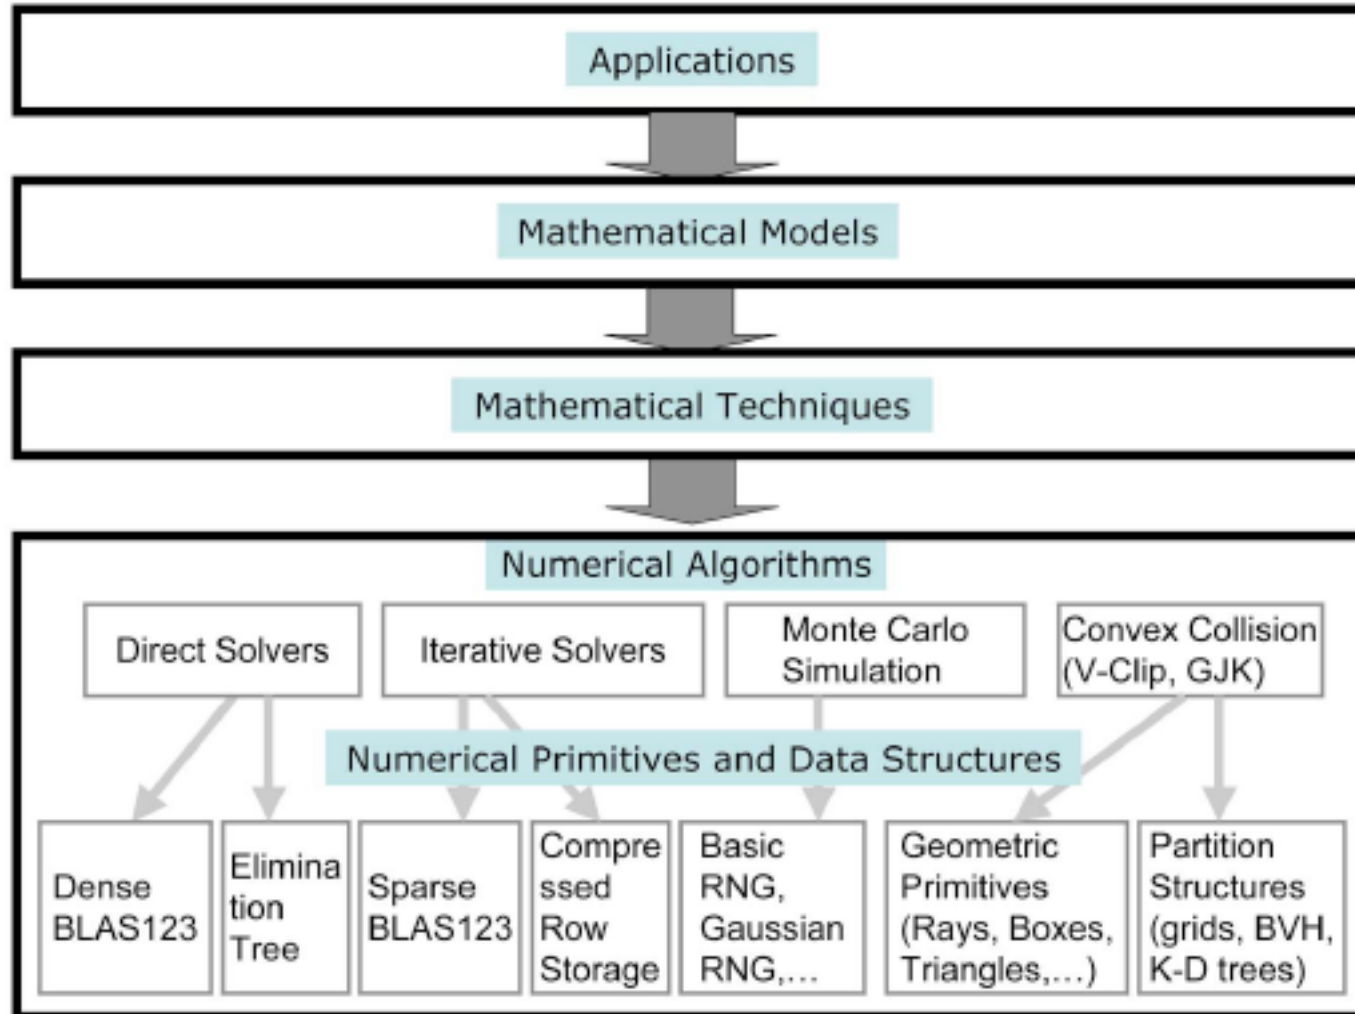
Synthesis



What if the tumor
progresses?


- Chen et al., *Convergence of Recognition, Mining, and Synthesis Workloads and Its Implications*. Proceedings of the IEEE, May 2008.

RMS App Commonality




Apps and Kernels Tower:

What are the problems?



- Old Conventional Wisdom: Use old programs to evaluate future computers
 - For example, SPEC2006, EEMBC
 - Tied to peculiarities of code artifact vs. fundamentals
 - Black-box benchmarks: don't understand or change internals
- Berkeley View
 - Computer HW and SW designers must *understand* applications
 - Killer apps for future systems are not yet known: understand the building blocks and algorithmic trends


Phillip Colella's "Seven dwarfs"



High-end simulation in the physical sciences = 7 numerical methods:


1. Structured Grids (including locally structured grids, e.g. Adaptive Mesh Refinement)
 2. Unstructured Grids
 3. Fast Fourier Transform
 4. Dense Linear Algebra
 5. Sparse Linear Algebra
 6. Particles
 7. Monte Carlo
- A dwarf is a pattern of computation and communication
 - Dwarfs are well-defined targets from algorithmic, software, and architecture standpoints

Do dwarfs work well outside HPC?




- Examine use of 7 dwarfs elsewhere
 1. Embedded Computing (EEMBC benchmark)
 2. Desktop/Server Computing (SPEC2006)
 3. Data Base / Text Mining Software
 - Advice from Jim Gray of Microsoft and Joe Hellerstein of UC
 4. Games/Graphics/Vision
 5. Machine Learning
 - Advice from Mike Jordan and Dan Klein of UC Berkeley
- Result: Added 7 more dwarfs, revised 2 original dwarfs, renumbered list

Dwarf Use (Red Important → Blue Not)




	Embed
1 Finite State Mach.	Red
2 Combinational	Red
3 Graph Traversal	Red
4 Structured Grid	Red
5 Dense Matrix	Red
6 Sparse Matrix	Yellow
7 Spectral (FFT)	Yellow
8 Dynamic Prog	Yellow

1. Embedded (42 EEMBC benchmarks)




Dwarf Use (Red Important → Blue Not)



Embed SPEC


- 1 Finite State Mach.
- 2 Combinational
- 3 Graph Traversal
- 4 Structured Grid
- 5 Dense Matrix
- 6 Sparse Matrix
- 7 Spectral (FFT)
- 8 Dynamic Prog
- 9 Particles
- 10 MapReduce



2. Desktop/Server (28 SPEC2006 benchmarks)




Dwarf Use (Red Important → Blue Not)



Embed SPEC DB

- 1 Finite State Mach.
- 2 Combinational
- 3 Graph Traversal
- 4 Structured Grid
- 5 Dense Matrix
- 6 Sparse Matrix
- 7 Spectral (FFT)
- 8 Dynamic Prog
- 9 Particles
- 10 MapReduce
- 11 Backtrack/ B&B
- 12 Graphical Models




3. Database / Text Mining

ORACLE

MySQL

Google

Dwarf Use (Red Important → Blue Not)




Embed SPEC DB Games

1 Finite State Mach.	Red	Red	Red	Yellow
2 Combinational	Red	Blue	Green	Blue
3 Graph Traversal	Red	Yellow	Yellow	Yellow
4 Structured Grid	Red	Red	Blue	Yellow
5 Dense Matrix	Red	Red	Yellow	Red
6 Sparse Matrix	Yellow	Yellow	Blue	Red
7 Spectral (FFT)	Yellow	Blue	Blue	Yellow
8 Dynamic Prog	Yellow	Blue	Red	Blue
9 Particles	Blue	Yellow	Blue	Yellow
10 MapReduce	Blue	Green	Red	Blue
11 Backtrack/ B&B	Blue	Blue	Yellow	Blue
12 Graphical Models	Blue	Blue	Yellow	Blue
13 Unstructured Grid	Blue	Blue	Blue	Yellow

4. Video Games



Dwarf Use (Red Important → Blue Not)



5. Machine Learning

	Embed	SPEC	DB	Games	ML
1 Finite State Mach.	Red	Red	Red	Yellow	Yellow
2 Combinational	Red	Blue	Green	Blue	Green
3 Graph Traversal	Red	Yellow	Yellow	Yellow	Red
4 Structured Grid	Red	Red	Blue	Yellow	Blue
5 Dense Matrix	Red	Red	Yellow	Red	Red
6 Sparse Matrix	Yellow	Yellow	Blue	Red	Red
7 Spectral (FFT)	Yellow	Blue	Blue	Yellow	Yellow
8 Dynamic Prog	Yellow	Blue	Red	Blue	Red
9 Particles	Blue	Yellow	Blue	Yellow	Blue
10 MapReduce	Blue	Green	Red	Blue	Red
11 Backtrack/ B&B	Blue	Blue	Yellow	Blue	Red
12 Graphical Models	Blue	Blue	Yellow	Blue	Red
13 Unstructured Grid	Blue	Blue	Blue	Yellow	Yellow






Embed SPEC DB Games ML HPC

	Map	Reduce	Shuffle	Sort	Group	Write
1 Finite State Mach.	Red	Yellow	Blue	Blue	Blue	Blue
2 Combinational	Red	Blue	Green	Blue	Green	Blue
3 Graph Traversal	Red	Yellow	Yellow	Red	Red	Blue
4 Structured Grid	Red	Blue	Yellow	Blue	Blue	Red
5 Dense Matrix	Red	Yellow	Red	Red	Red	Red
6 Sparse Matrix	Yellow	Blue	Red	Red	Red	Red
7 Spectral (FFT)	Yellow	Blue	Yellow	Yellow	Yellow	Red
8 Dynamic Prog	Yellow	Blue	Red	Blue	Red	Blue
9 Particles	Blue	Yellow	Blue	Yellow	Blue	Red
10 MapReduce	Blue	Green	Red	Blue	Red	Red
11 Backtrack/ B&B	Blue	Yellow	Blue	Red	Red	Blue
12 Graphical Models	Blue	Yellow	Blue	Red	Red	Blue
13 Unstructured Grid	Blue	Yellow	Yellow	Yellow	Yellow	Red

Roles of Dwarfs


1. Give us a vocabulary/organization to talk across disciplinary boundaries
2. Define minimum set of necessary functionality for new hardware/software systems
3. Define building blocks for creating libraries that cut across application domains
4. “Anti-benchmarks” not tied to code or language artifacts \Rightarrow encourage innovation in algorithms, languages, data structures, and/or hardware
5. They decouple research, allowing analysis of HW & SW programming support without waiting years for full app development



Hardware Tower:

What are the problems?


- Power limits leading edge chip designs
 - Intel Tejas Pentium 4 cancelled due to power issues
- Yield on leading edge processes dropping dramatically
 - IBM quotes yields of 10–20% on 8-processor Cell
 - PS3: only 6 Cell SPUs available to programmer
- Design/validation leading edge chip is becoming unmanageable
 - Verification teams > design teams on leading edge processors



HW Solution: Small is Beautiful


- Expect modestly pipelined (5- to 9-stage) CPUs, FPUs, vector, SIMD PEs
 - Small cores not much slower than large cores
- Parallel is energy efficient path to performance:
 $\text{Power} = CV^2F$
 - Lower voltage, and increase parallelism lowers energy per op
- Redundant processors can improve chip yield
 - Cisco Metro 188 CPUs + 4 spares; Sun Niagara sells 6 or 8 CPUs; BlueGene CPU has 1 spare
- Small, regular processors easier to verify
- One size fits all?
 - Amdahl's Law \Rightarrow Heterogeneous processors?

Number of Cores/Socket



- We need revolution, not evolution
- Software or architecture alone can't fix parallel programming problem, need innovations in both
- “Multicore” 2X cores per generation: 2, 4, 8, ...
- “Manycore” 100s is highest performance per unit area, and per Watt, then 2X per generation: 64, 128, 256, 512, 1024 ...
- **Multicore architectures, programming models, and applications good for 2 to 32 cores won't evolve to Manycore systems of 1000's of cores
⇒ Desperately need HW/SW models that work for Manycore or will run out of steam
(as ILP ran out of steam at 4 instructions)**

7 Questions for Parallelism



- **Applications:**

1. What are the apps?
2. What are kernels of apps?

- **Hardware:**


3. What are the HW building blocks?
4. How to connect them?

- **Programming Model & Systems Software:**


5. How to describe apps and kernels?
6. How to program the HW?

- **Evaluation:**

7. How to measure success?



(Inspired by a view of the Golden Gate Bridge from Berkeley)



Programming Model: What are the problems?

- Primary recent focus on correctness, not performance
 - Relied on “Moore’s Law” to make programs faster
- New generation of performance programmers
 - Why parallel if performance doesn’t matter?
- Programming model must balance *productivity* and *implementation efficiency*
 - Enable software industry for manycore
 - Usable by most programmers
 - NSA conference: “C” programmers

Old CW in Programming Models

- Design new hardware with exotic performance features
- Hand to software communication
 - Develop new compiler technology and wait for maturity and integration into commercial compilers
 - Takes ~10 years in practice
- Compilers should
 - Compiler arbitrary code efficiently
 - Hide performance issues from programmer
 - Run quickly (secs-mins, human is in the loop)
- Search for the holy grail language
 - One language for all problems

New CW in Programming Models

- Feature creep in languages annotations
→ de facto new languages
- Programs written in many languages
 - Python, C++, Perl, Java, Javascript, C#,...
- Automatic performance tuning
 - Use machine time in place of human time for tuning
 - Search over possible implementations
 - Autotuned libraries for dwarfs (up to 10x speedup)
 - Spectral (FFTW, Spiral) •Sparse (OSKI)
 - Dense (Atlas, PHiPAC) •Structured (OSKI')



Measuring Success: What are the problems?

1. \approx Only companies can build HW, and it takes years
2. Software people don't start working hard until hardware arrives
 - 3 months after HW arrives, SW people list everything that must be fixed, then we all wait 4 years for next iteration of HW/SW
3. How get 1000 CPU systems in hands of researchers to innovate in timely fashion on in algorithms, compilers, languages, OS, architectures, ... ?
4. Can avoid waiting years between HW/SW iterations?

Rapid Prototyping with FPGAs



RAMP Blue, January 2007


- 256 RISC cores @100MHz
- Works! Runs UPC version of NAS benchmarks.

- “Research Accelerator for Multi-Processors”
 - Multi-University collaboration developing FPGA “gateway” for manycore emulations (10 faculty at UCB, CMU, MIT, Stanford, Texas, Washington)
- Enables rapid interaction between hardware and software developers
 - “Tapeout” every day, not once in five years
 - Fast enough (100MHz) for software development
- RAMP Design Language (RDL) provides “gateway linker” and **cycle-accurate timing models**.
 - *All operations (DRAM access, FP multiply, disk access) take exact same number of clock cycles as on desired target machine*
- Multiple machine styles in progress
 - RAMP Blue (UC Berkeley) cluster/message-passing
 - RAMP Red (Stanford) transactional memory
 - RAMP White (Everyone) cache-coherent CMP

NRC Report Recommendations

- Invest in research in and development of **algorithms that can exploit parallel processing**.
- Invest in research in and development of **programming methods that will enable efficient use of parallel systems** not only by parallel-systems experts but also by typical programmers.
- Focus long-term efforts on **rethinking of the canonical computing** “stack” applications, programming language, compiler, runtime, virtual machine, operating system, hypervisor, and architecture—in light of parallelism and resource-management challenges.
- Invest in **research on and development of parallel architectures driven by applications**, including enhancements of chip multiprocessor systems and conventional data-parallel architectures, cost-effective designs for application-specific architectures, and support for radically different approaches.


Today's Supercomputers and the Brain




compute only

Transistors	Memory	Clock (GHz)	Power (W)	Weight (kg)	Size (📏📏)
9×10^{12}	2×10^{15}	2	1.2×10^6	18,800	36,000

9×10^{10}	1×10^{15}	$10^{-9} - 10^{-5}$	20	1.5	
Neurons	Syn. Conn.				




Real time simulation at the exascale



Ananathanarayanan et al., "The Cat is out of the Bag: Cortical Simulations with 10^9 Neurons and 10^{15} Synapses", Proceedings of SC09.

Recent Evidence for Petabyte Size Memory

- Bartol et al., “Nanoconnectomic Upper Bound on the Variability of Synaptic Plasticity”, Jan. 2016, eLife.
- 26 sizes of synapses corresponds to about 4.7 bits per synapse and thus about 4 – 5 Petabytes



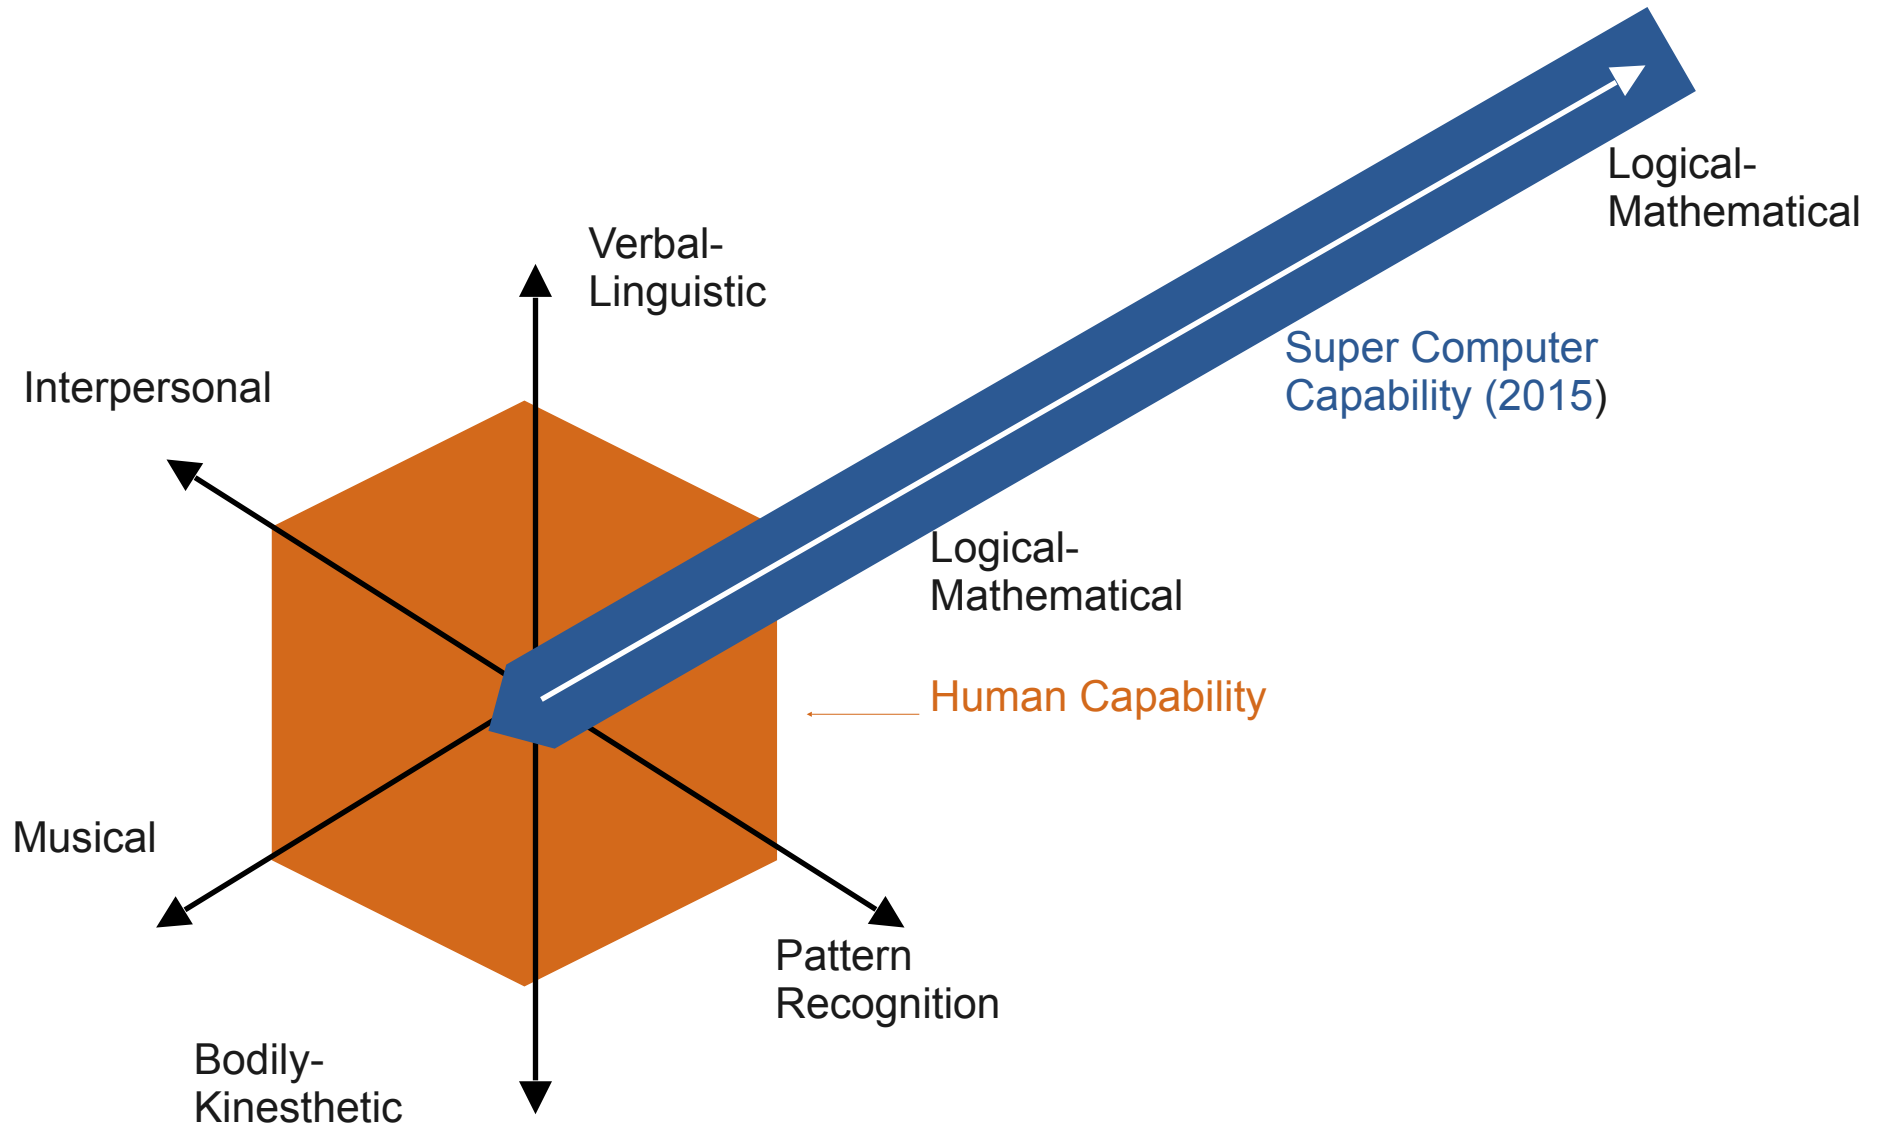
from <http://www.salk.edu/news-release/memory-capacity-of-brain-is-10-times-more-than-previously-thought/>


Compute Power of the Human Brain

- Estimate of compute power for the human brain is about 1-10 Exaflops and 4-5 Petabytes
- Three different paths lead to about the same estimate
- A digital computer with this performance might be available in about 2024 with a power consumption of at best 20–30 MW (goal of the Exascale project)
- The human brain takes 20 W
- A digital exaflops computer using CMOS technology will still be a factor of a million away from brain power




Current State of Supercomputers





From A. Merolla
et al., *Science*,
Aug. 8, 2014




NRC Report Recommendations

- Invest in research and development to make computer systems **more power-efficient** at all levels of the system, including software, application-specific approaches, and alternative devices.
- To promote cooperation and innovation by sharing encourage **development of open interface standards for parallel programming** rather than proliferating proprietary programming environments.
- Invest in the development of **tools and methods to transform legacy applications to parallel systems**.
- Incorporate in **computer science education an increased emphasis on parallelism**, and use a variety of methods and approaches to prepare students better for the types of computing resources that they will encounter in their careers.

Amdahl's Law

- Speedup due to enhancement E :

$$\text{Speedup}(E) = \frac{\text{Execution Time without } E}{\text{Execution Time with } E} = \frac{\text{Performance with } E}{\text{Performance without } E}$$



- Suppose that enhancement E accelerates a fraction F of the task by a factor S and the remainder of the task is unaffected:

$$\text{Execution time (with } E) = ((1 - F) + F/S) \cdot \text{Execution time (without } E)$$

$$\text{Speedup (with } E) = \frac{1}{(1 - F) + F/S}$$

- Design Principle: Make the common case fast!

Why EEC 171?

- Old CW: Don't bother parallelizing your application, as you can just wait a little while and run it on a much faster sequential computer.
- New CW: It will be a very long wait for a faster sequential computer.
- Old CW: Increasing clock frequency is the primary method of improving processor performance.
- New CW: Increasing parallelism is the primary method of improving processor performance.
- Old CW: Less than linear scaling for a multiprocessor application is failure.
- New CW: Given the switch to parallel computing, any speedup via parallelism is a success.

Extracting Yet More Performance

- Two options:
 - Increase the depth of the pipeline to increase the clock rate — superpipelining
 - How does this help performance? (What does it impact in the performance equation?)
 - Fetch (and execute) more than one instruction at one time (expand every pipeline stage to accommodate multiple instructions) — multiple-issue
 - How does this help performance? (What does it impact in the performance equation?)
- Today's topic!
$$\frac{\text{seconds}}{\text{program}} = \frac{\text{instructions}}{\text{program}} \times \frac{\text{cycles}}{\text{instruction}} \times \frac{\text{seconds}}{\text{cycle}}$$

Instruction vs Machine Parallelism

- Instruction-level parallelism (ILP) of a program—a measure of the average number of instructions in a program that a processor might be able to execute at the same time
- Mostly determined by the number of true (data) dependencies and procedural (control) dependencies in relation to the number of other instructions
- ILP is traditionally “extracting parallelism from a single instruction stream working on a single stream of data”

Instruction vs Machine Parallelism

- Machine parallelism of a processor—a measure of the ability of the processor to take advantage of the ILP of the program
- Determined by the number of instructions that can be fetched and executed at the same time
- ***To achieve high performance, need both ILP and machine parallelism***

Why is ILP a good idea? If you were designing a computer system, why would you choose ILP instead of, say, multiple processors?

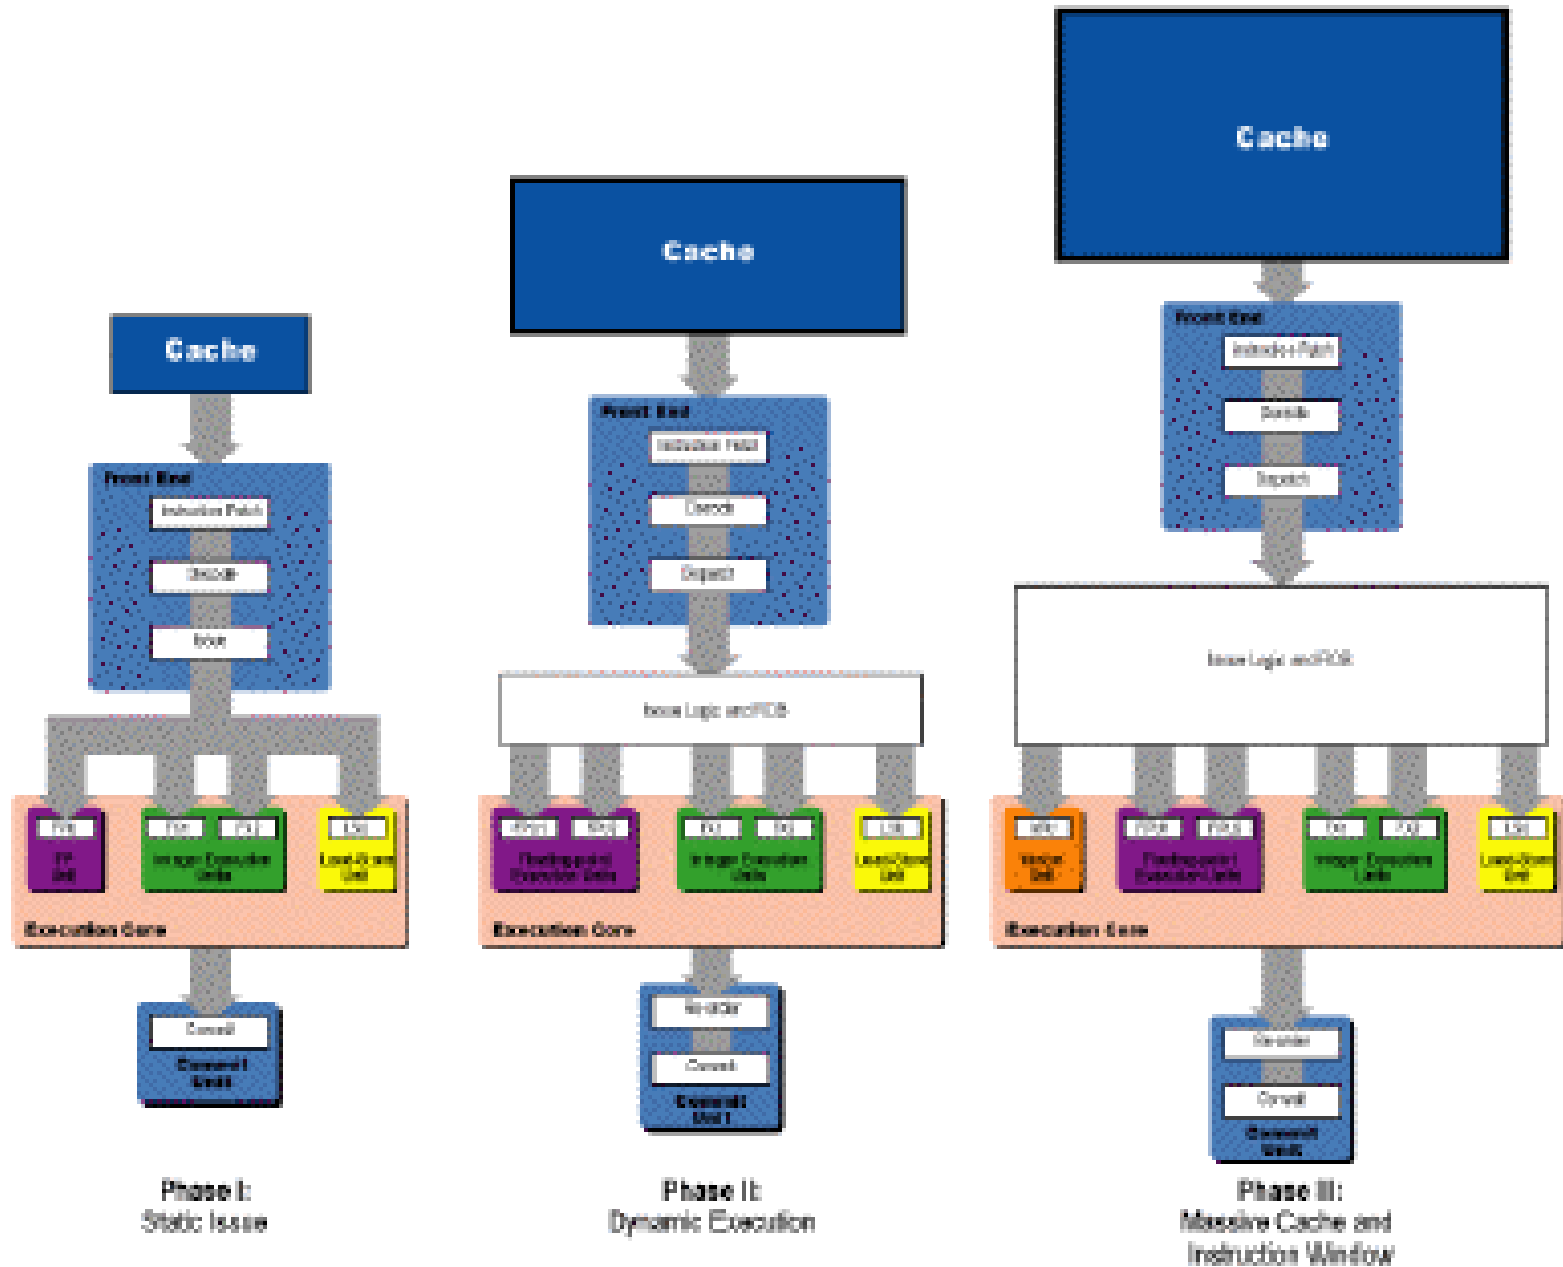
What kind of code has lots of ILP?

What kind of code has little ILP?

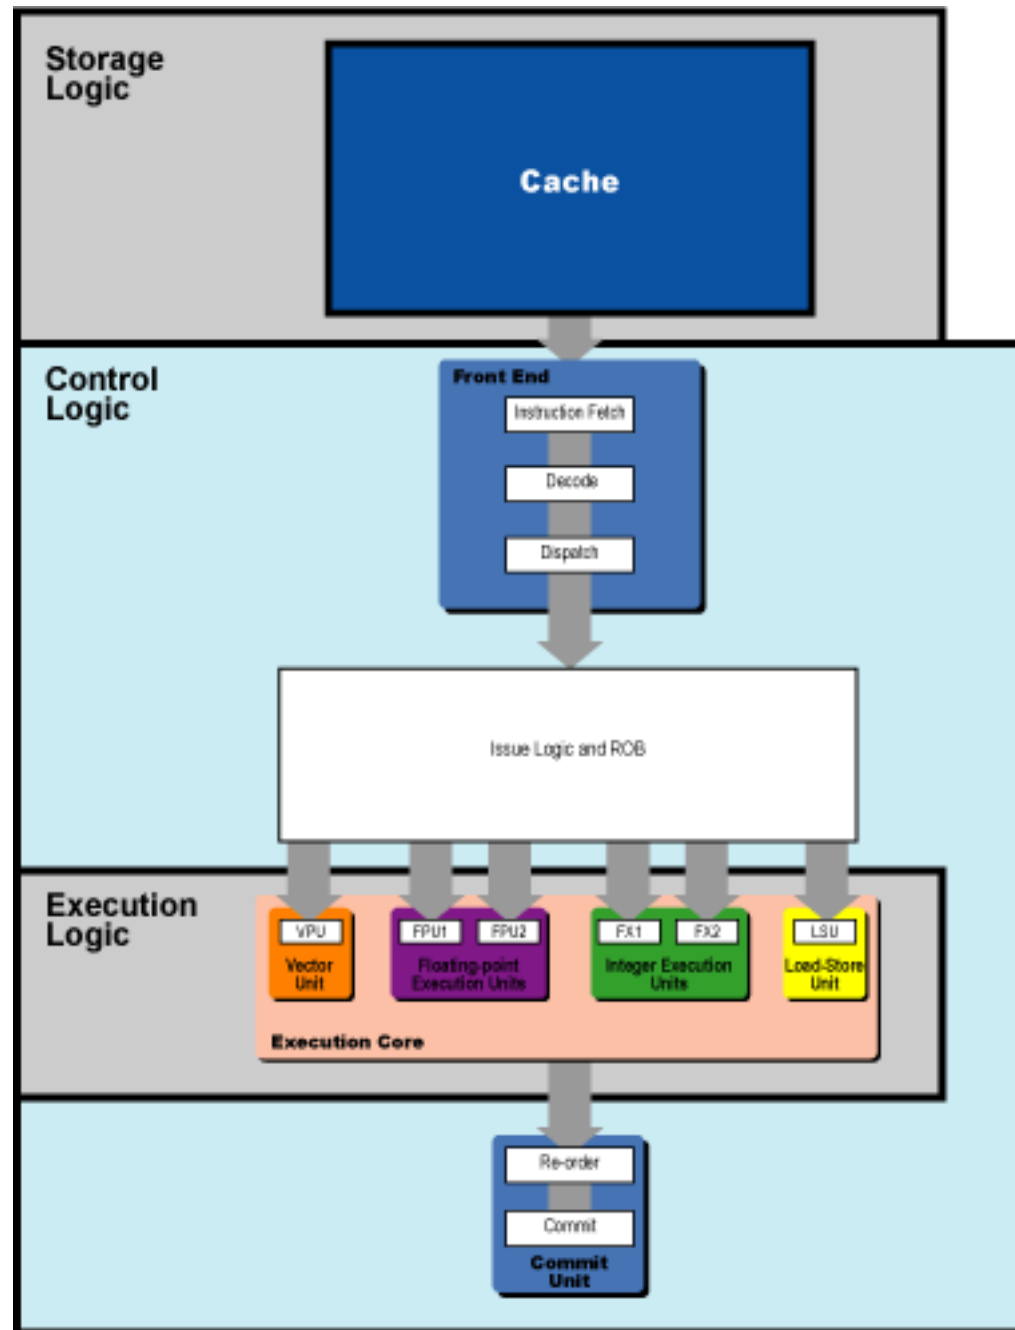
Machine Parallelism

- There are 2 main approaches for machine parallelism. Responsibility of resolving hazards is ...
 - Primarily hardware-based—“dynamic issue”, “superscalar”
 - Primarily software-based—“VLIW”


Growing complexity ...



Small fraction for datapath



AMD “Deerhound” (K8L)



How does out-of-order issue help?

How does out-of-order completion
help?

How does register renaming help?

Static Multiple Issue Machines (VLIW)

- Static multiple-issue processors (aka VLIW) use the compiler to decide which instructions to issue and execute simultaneously
- Issue packet—the set of instructions that are bundled together and issued in one clock cycle—think of it as one large instruction with multiple operations
- The mix of instructions in the packet (bundle) is usually restricted—a single “instruction” with several predefined fields
- The compiler does static branch prediction and code scheduling to reduce (ctrl) or eliminate (data) hazards

What's good about VLIW?

What's bad about VLIW?

Predication

- Predication can be used to eliminate branches by making the execution of an instruction dependent on a “predicate”, e.g.,

```
if (p) { statement 1 } else { statement 2 }
```

would normally compile using two branches. (Why?) With predication it would compile as

```
(p) statement 1
```

```
(~p) statement 2
```

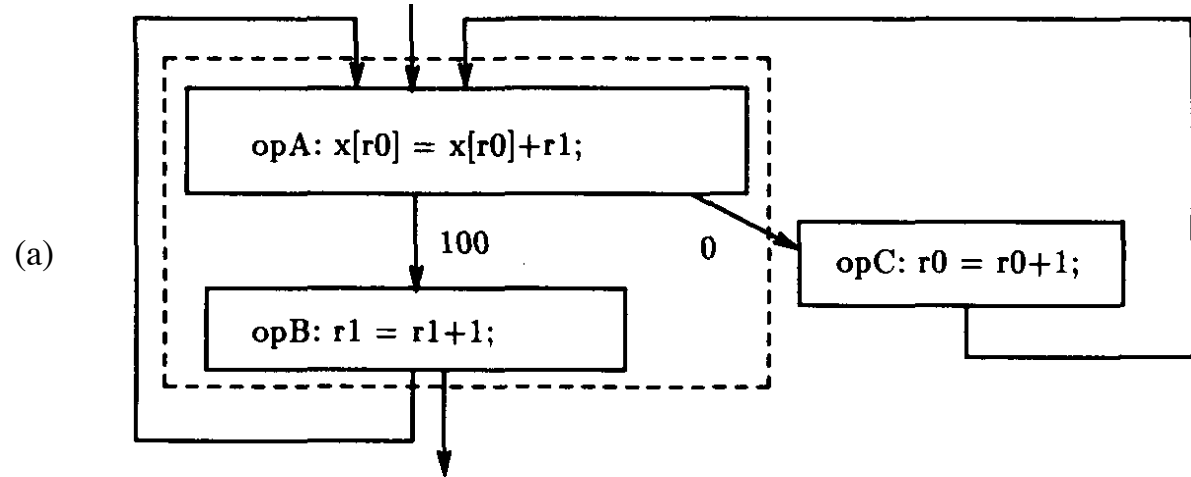
- The use of `(condition)` indicates that the instruction is committed only if `condition` is true
- Predication can be used to speculate as well as to eliminate branches

Speculation

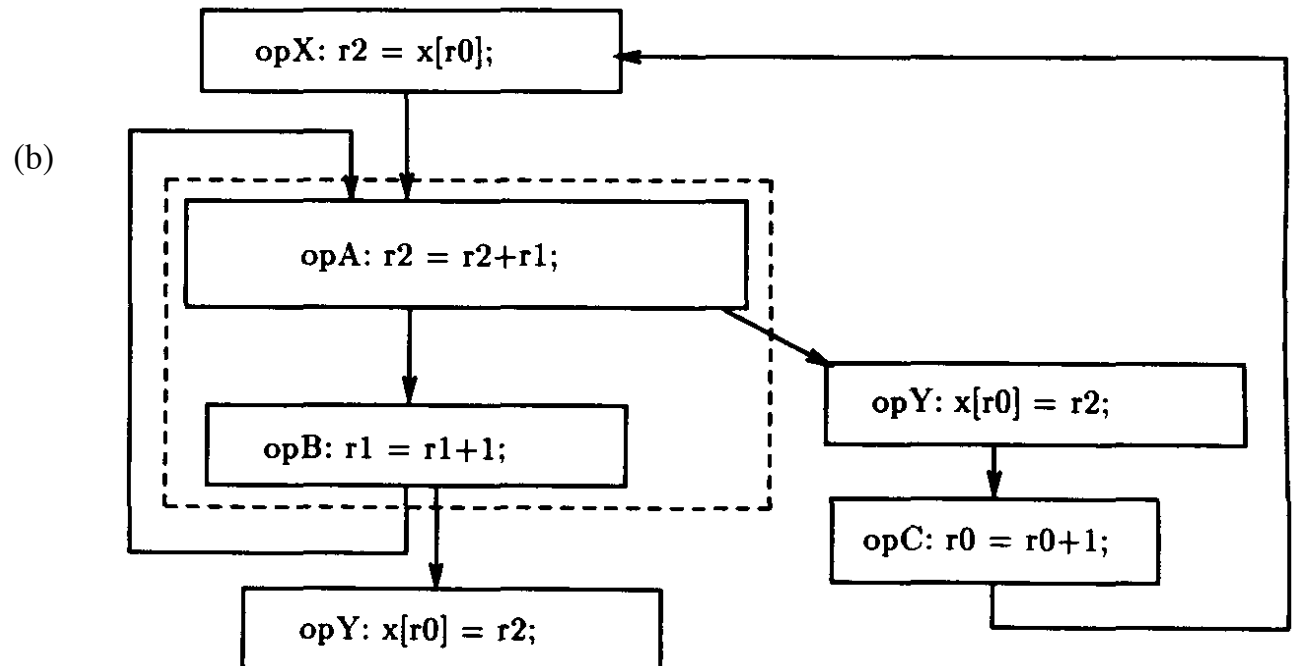
- Speculation is used to allow execution of future instructions that (may) depend on the speculated instruction
- Speculate on the outcome of a conditional branch (branch prediction)
 - Compare to out-of-order machine with branch prediction
- Speculate that a store (for which we don't yet know the address) that precedes a load does not refer to the same address, allowing the load to be scheduled before the store (load speculation)

Trace Scheduling

- Original loop allows us to increment either r_0 or r_1 :
 $x[r_0] = x[r_0] + r_1$



- Profile says incrementing r_1 is much more common
- Optimize for that case



ILP Summary

- Leverage Implicit Parallelism for Performance: Instruction Level Parallelism
- Loop unrolling by compiler to increase ILP
- Branch prediction to increase ILP
- Dynamic HW exploiting ILP
 - Works when can't know dependence at compile time
 - Can hide L1 cache misses
 - Code for one machine runs well on another

Flynn's Classification Scheme



- SISD – single instruction, single data stream
 - Uniprocessors
- SIMD – single instruction, multiple data streams
 - single control unit broadcasting operations to multiple datapaths
- MISD – multiple instruction, single data
 - no such machine (although some people put vector machines in this category)
- MIMD – multiple instructions, multiple data streams
 - aka multiprocessors (SMPs, MPPs, clusters, NOWs)

Continuum of Granularity

- “Coarse”

- Each processor is more powerful
- Usually fewer processors
- Communication is more expensive between processors
- Processors are more loosely coupled
- Tend toward MIMD


- “Fine”

- Each processor is less powerful
- Usually more processors
- Communication is cheaper between processors
- Processors are more tightly coupled
- Tend toward SIMD


What kind of problems are good
for coarse-grained parallelism?

What kind of problems are good
for fine-grained parallelism?

Simultaneous multithreading (SMT)




Centralized vs. Distributed Memory




2 Classes of Cache Coherence Protocols

- **Directory based** — Sharing status of a block of physical memory is kept in just one location, the directory
- **Snooping** — Every cache with a copy of data also has a copy of sharing status of block, but no centralized state is kept
- All caches are accessible via some broadcast medium (a bus or switch)
- All cache controllers monitor or snoop on the medium to determine whether or not they have a copy of a block that is requested on a bus or switch access


Types of Communication




1 to 1



1 to N



N to 1




N to M


Network Topology

Distributed Switched (Direct) Networks


2D **mesh** or grid of 16 nodes



2D **torus** of 16 nodes




hypercube of 16 nodes
($16 = 2^4$, so $n = 4$)




**Network
Bisection**

\leq full bisection bandwidth!

SIMD Instructions

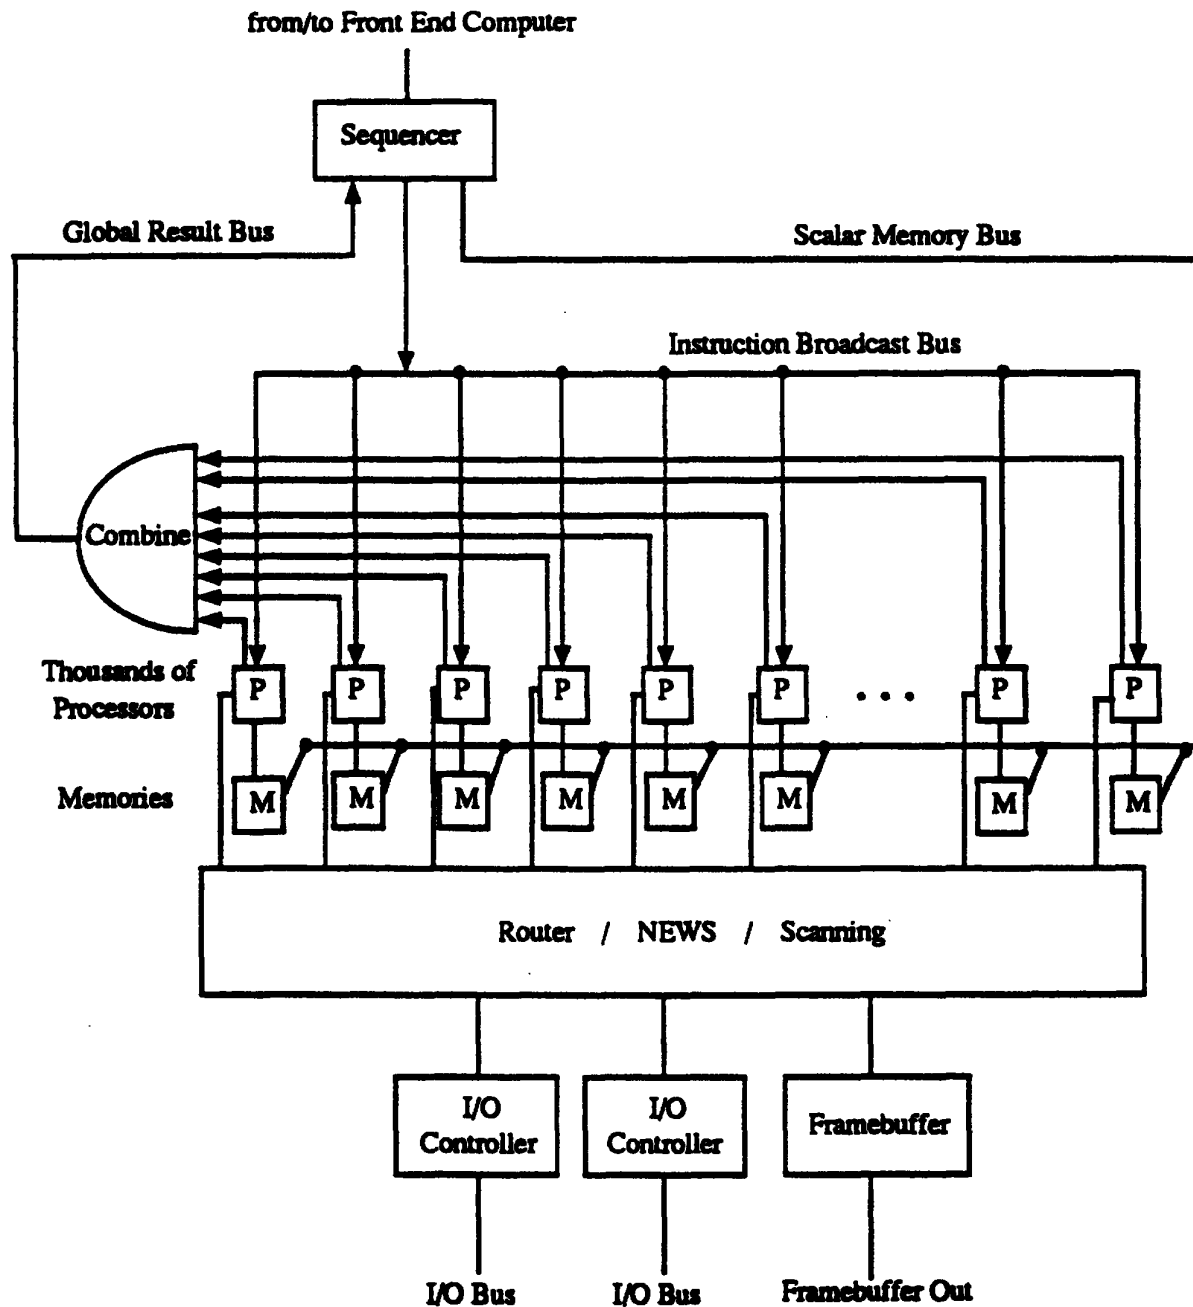


Cray-1 (1976)

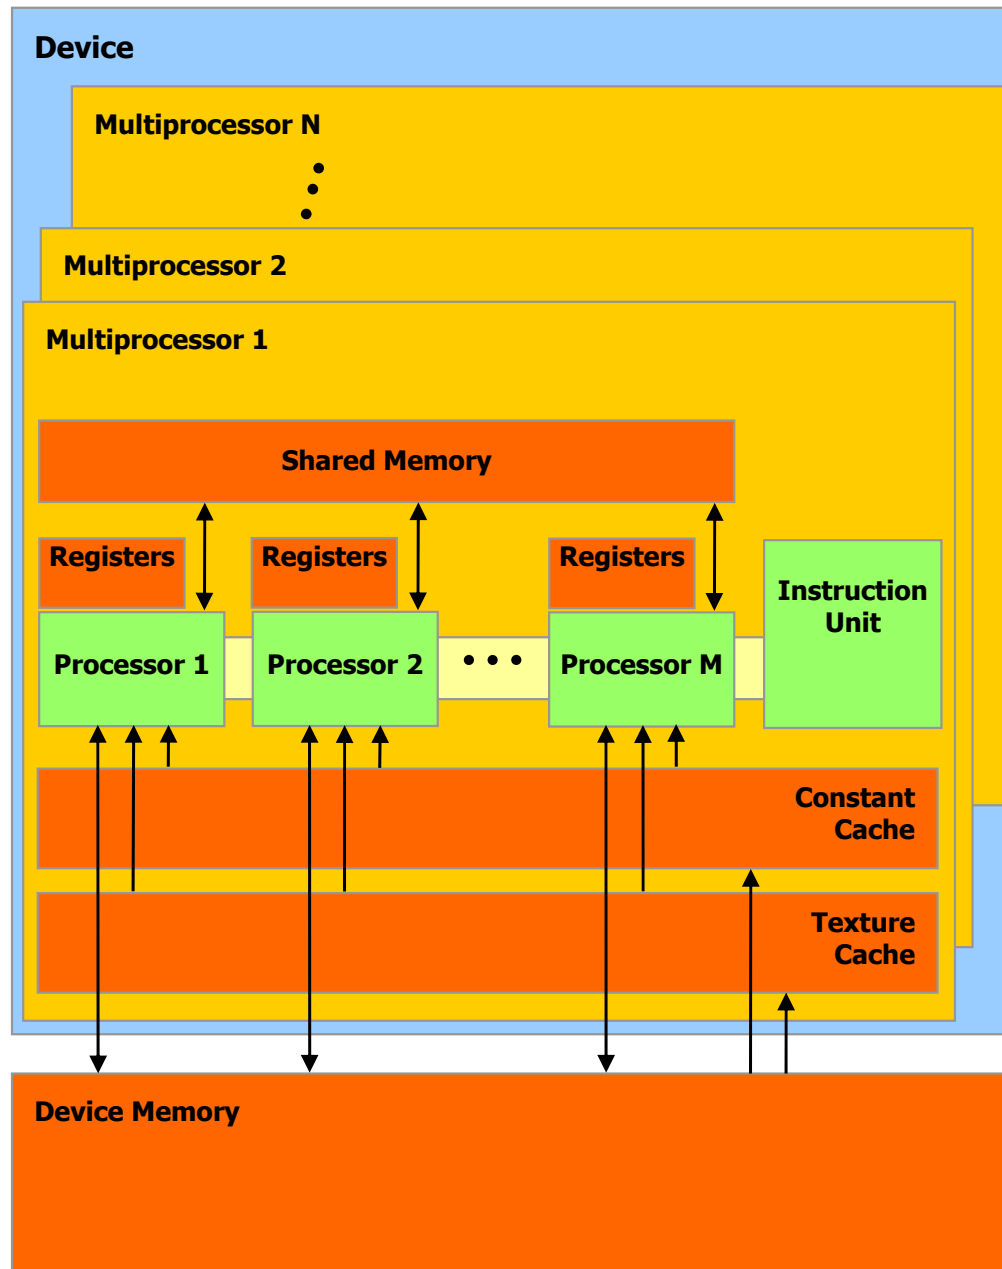


memory bank cycle 50 ns processor cycle 12.5 ns (80 MHz)

CM-2 Hardware Overview



CUDA Hardware Abstraction



Data-Parallel Algorithms

- Efficient algorithms require efficient building blocks
- Data-parallel building blocks
 - Map
 - Gather & Scatter
 - Reduce
 - Scan
 - Sort

Thank you all for a terrific quarter!

Good luck on your exams!